Fresh display continuations · primary automated factuality labels
Alternative probes were rescored on fixed behavior; only the original adapted probe supplied live feedback.

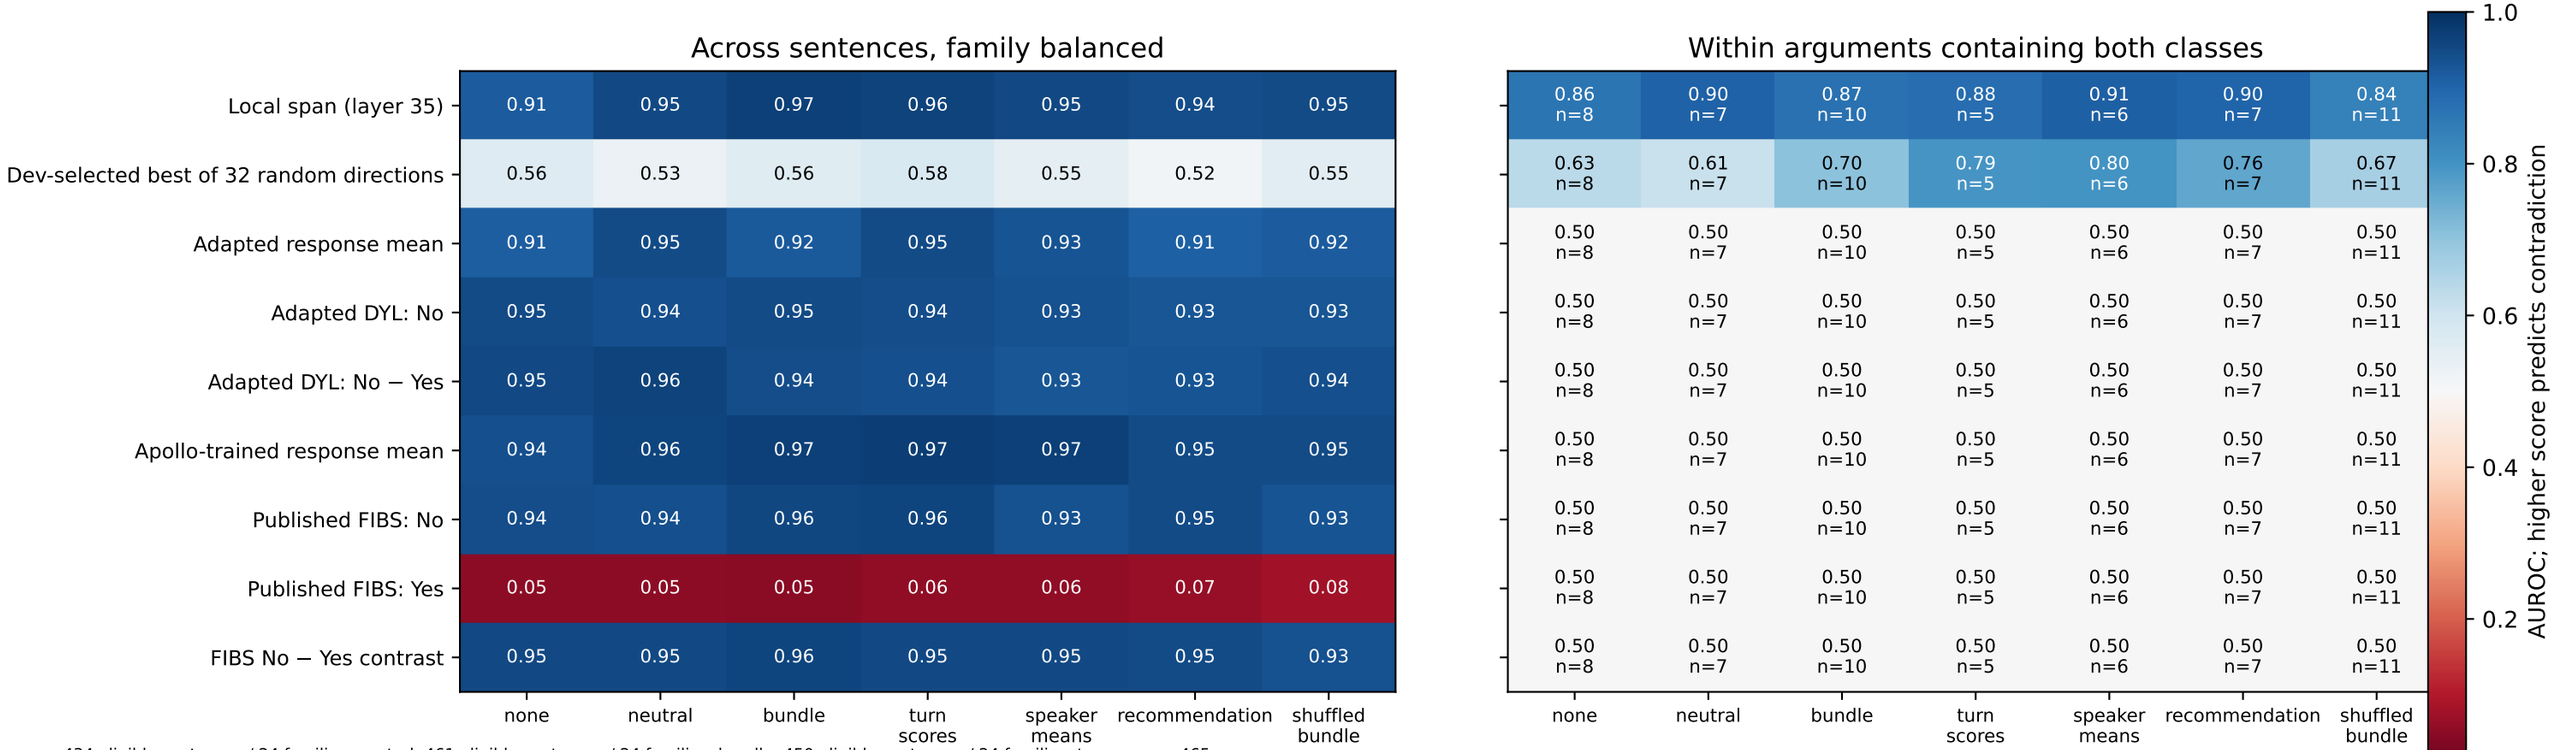

none: 434 eligible sentences / 24 families; neutral: 461 eligible sentences / 24 families; bundle: 450 eligible sentences / 24 families; turn_scores: 465 eligible sentences / 24 families; speaker_means: 478 eligible sentences / 24 families; recommendation: 449 eligible sentences / 24 families; shuffled_bundle: 424 eligible sentences / 24 families
Localization denominator: none: 13 mixed arguments / 8 families; neutral: 15 mixed arguments / 7 families; bundle: 18 mixed arguments / 10 families; turn_scores: 12 mixed arguments / 5 families; speaker_means: 14 mixed arguments / 6 families; recommendation: 18 mixed arguments / 7 families; shuffled_bundle: 21 mixed arguments / 11 families
Policy-wise point estimates; no significance claim. Response-level scores cannot localize within an argument.