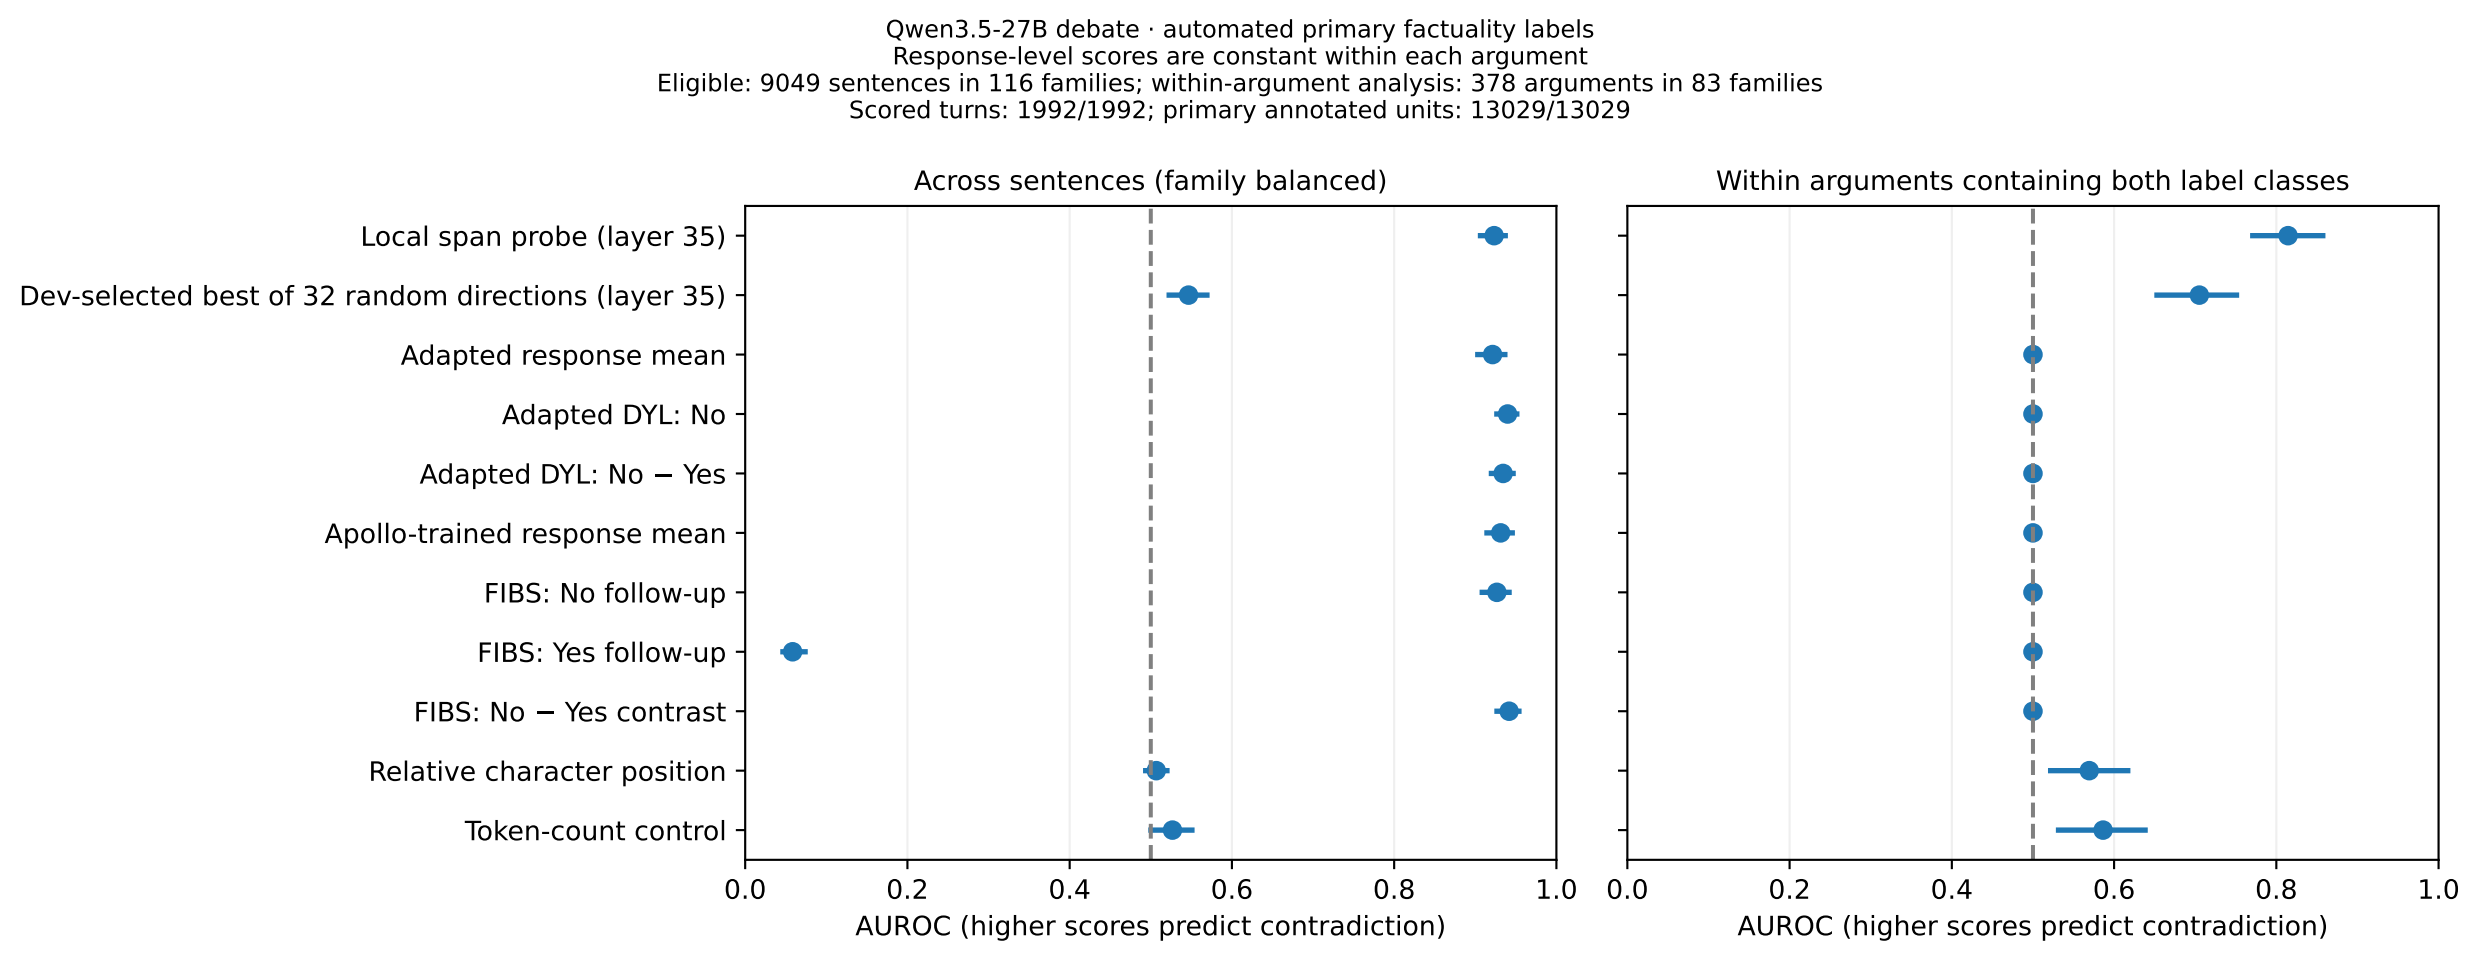
Qwen3.5-27B debate · automated primary factuality labels
Response-level scores are constant within each argument
Eligible: 9049 sentences in 116 families; within-argument analysis: 378 arguments in 83 families
Scored turns: 1992/1992; primary annotated units: 13029/13029
Across sentences (family balanced)
Within arguments containing both label classes
Local span probe (layer 35)
Dev-selected best of 32 random directions (layer 35)
Adapted response mean
Adapted DYL: No
Adapted DYL: No − Yes
Apollo-trained response mean
FIBS: No follow-up
FIBS: Yes follow-up
FIBS: No − Yes contrast
Relative character position
Token-count control
0.0
0.2
0.4
0.6
0.8
1.0
AUROC (higher scores predict contradiction)