Probe assistance depends on the judge

Prespecified primary comparisons · both Qwen 27B debater models pooled

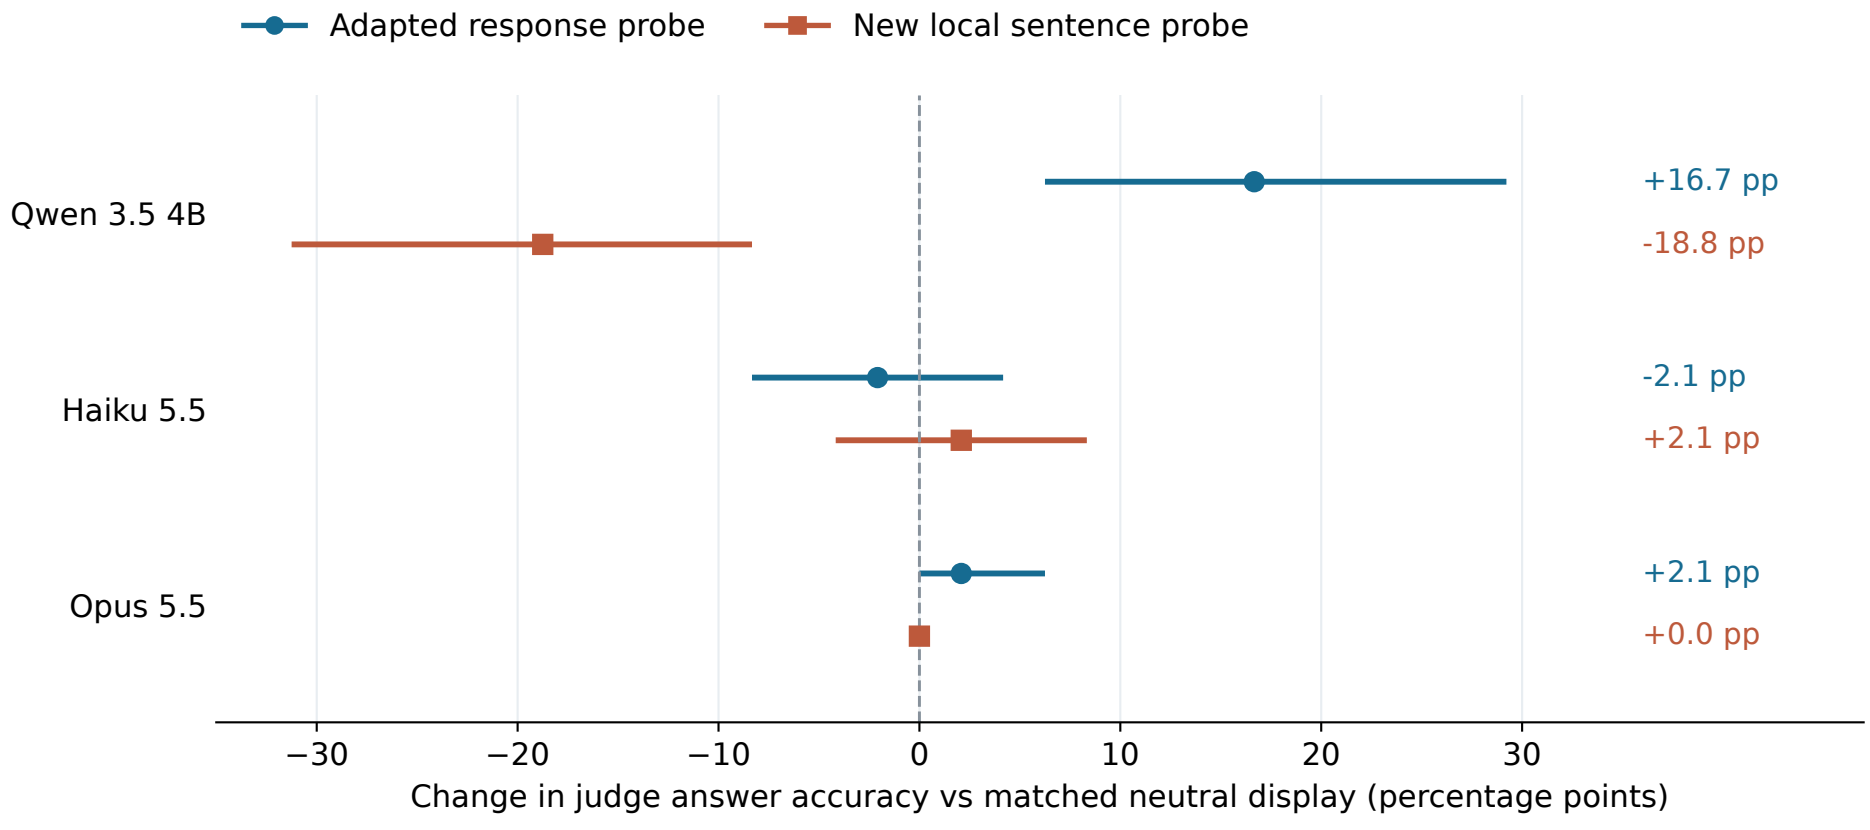

48 debate presentations from 24 original story families per point. Bars: unadjusted 95% family bootstrap intervals.
Positive values improve judge accuracy; negative values reduce it. These are paired effects, not raw baseline accuracy.