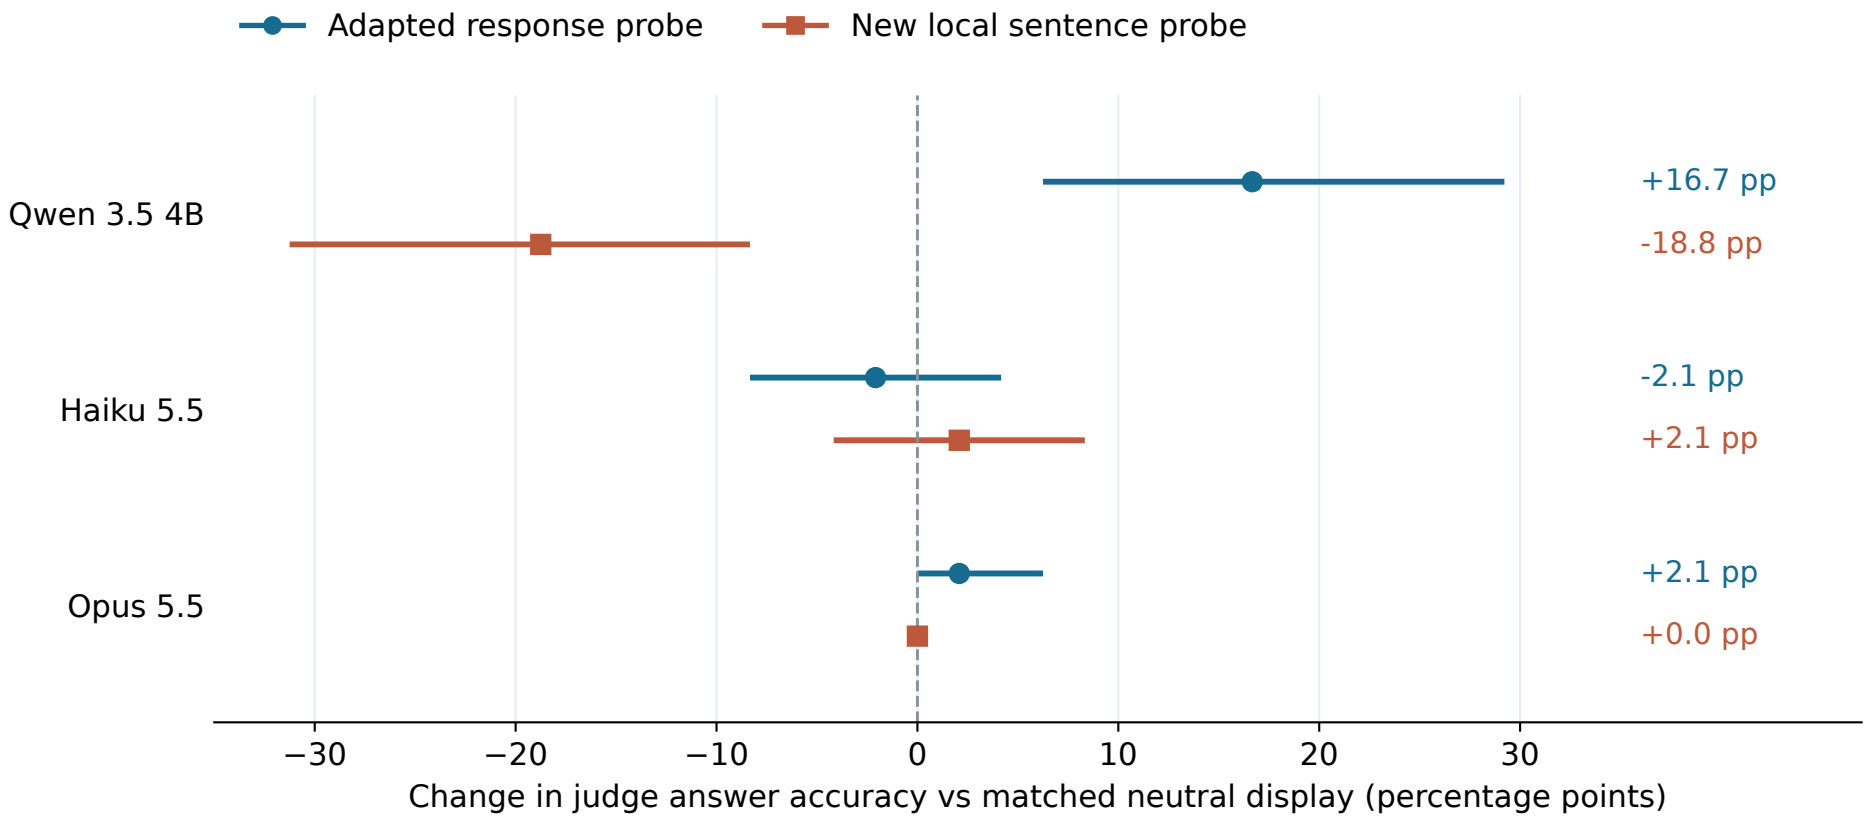
Probe assistance depends on the judge

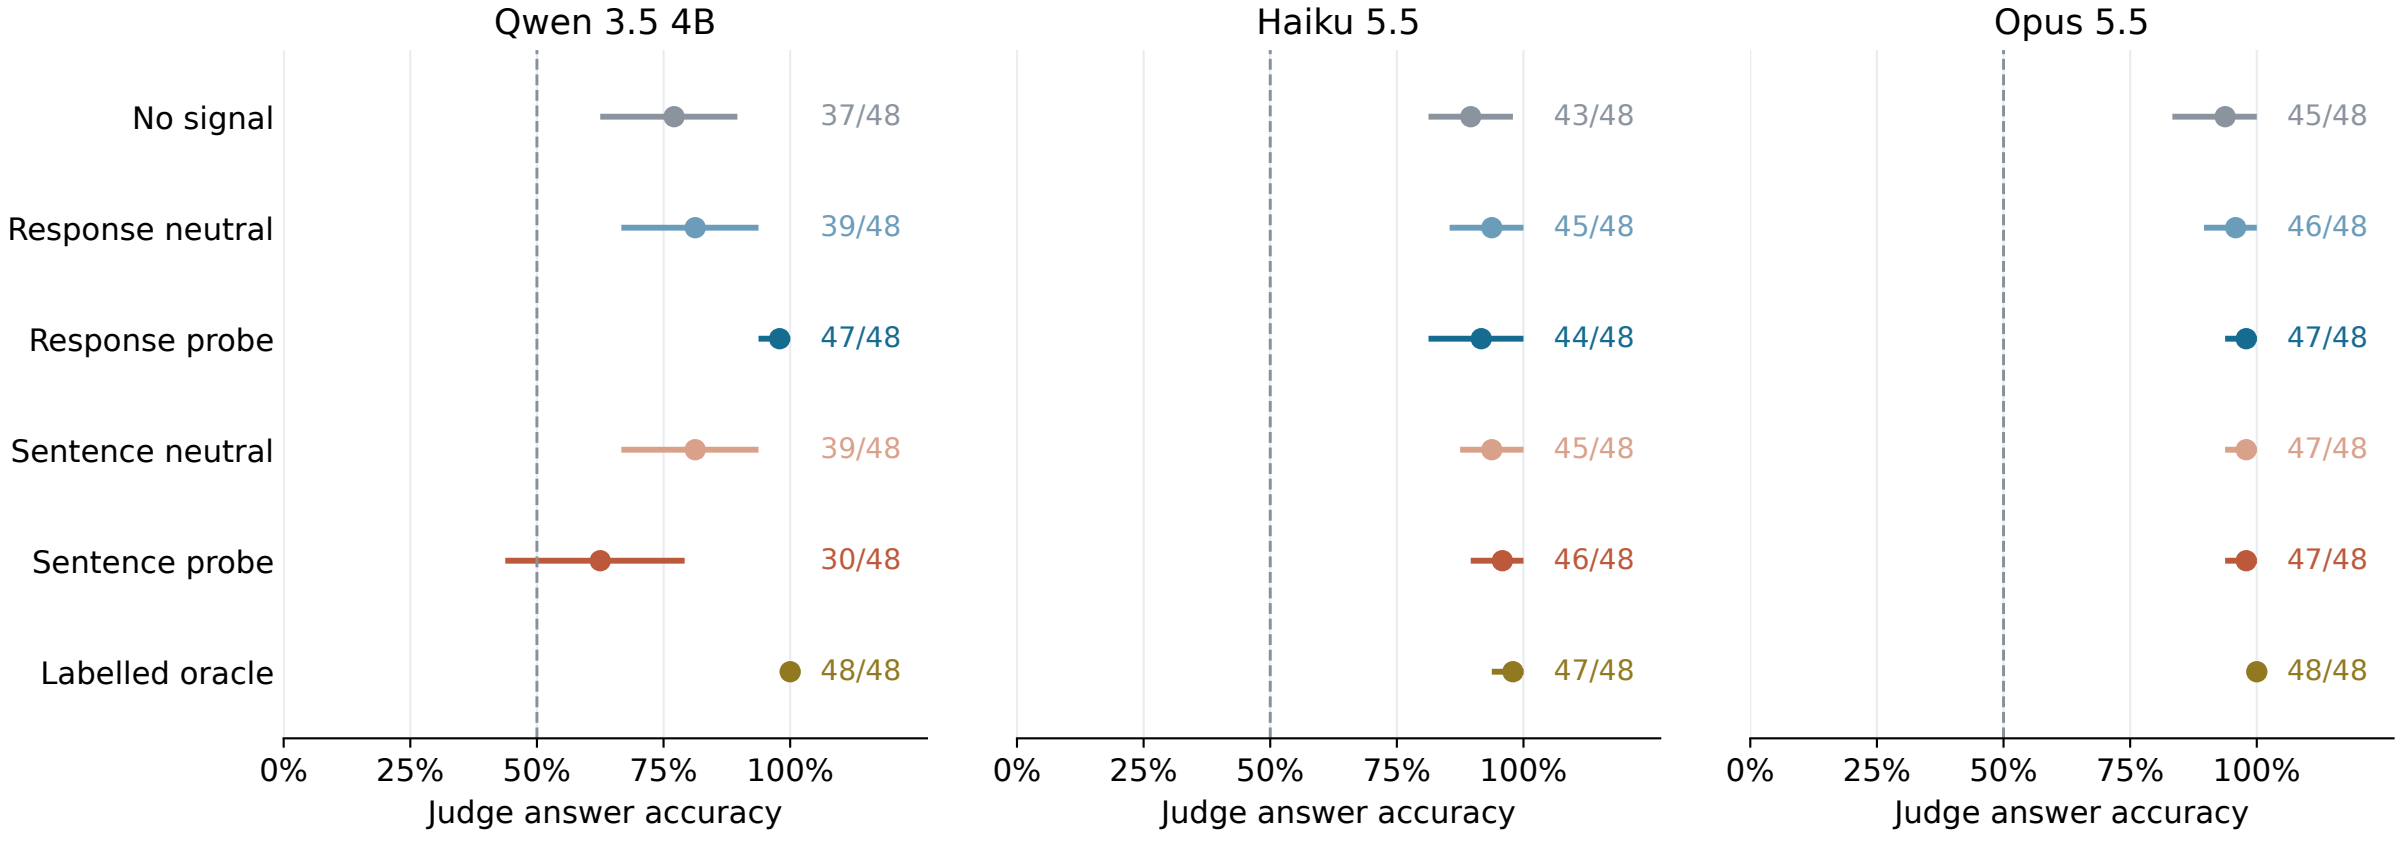
Prespecified primary comparisons · both Qwen 27B debater models pooled



48 debate presentations from 24 original story families per point. Bars: unadjusted 95% family bootstrap intervals.
Positive values improve judge accuracy; negative values reduce it. These are paired effects, not raw baseline accuracy.

Stronger judges start closer to the ceiling

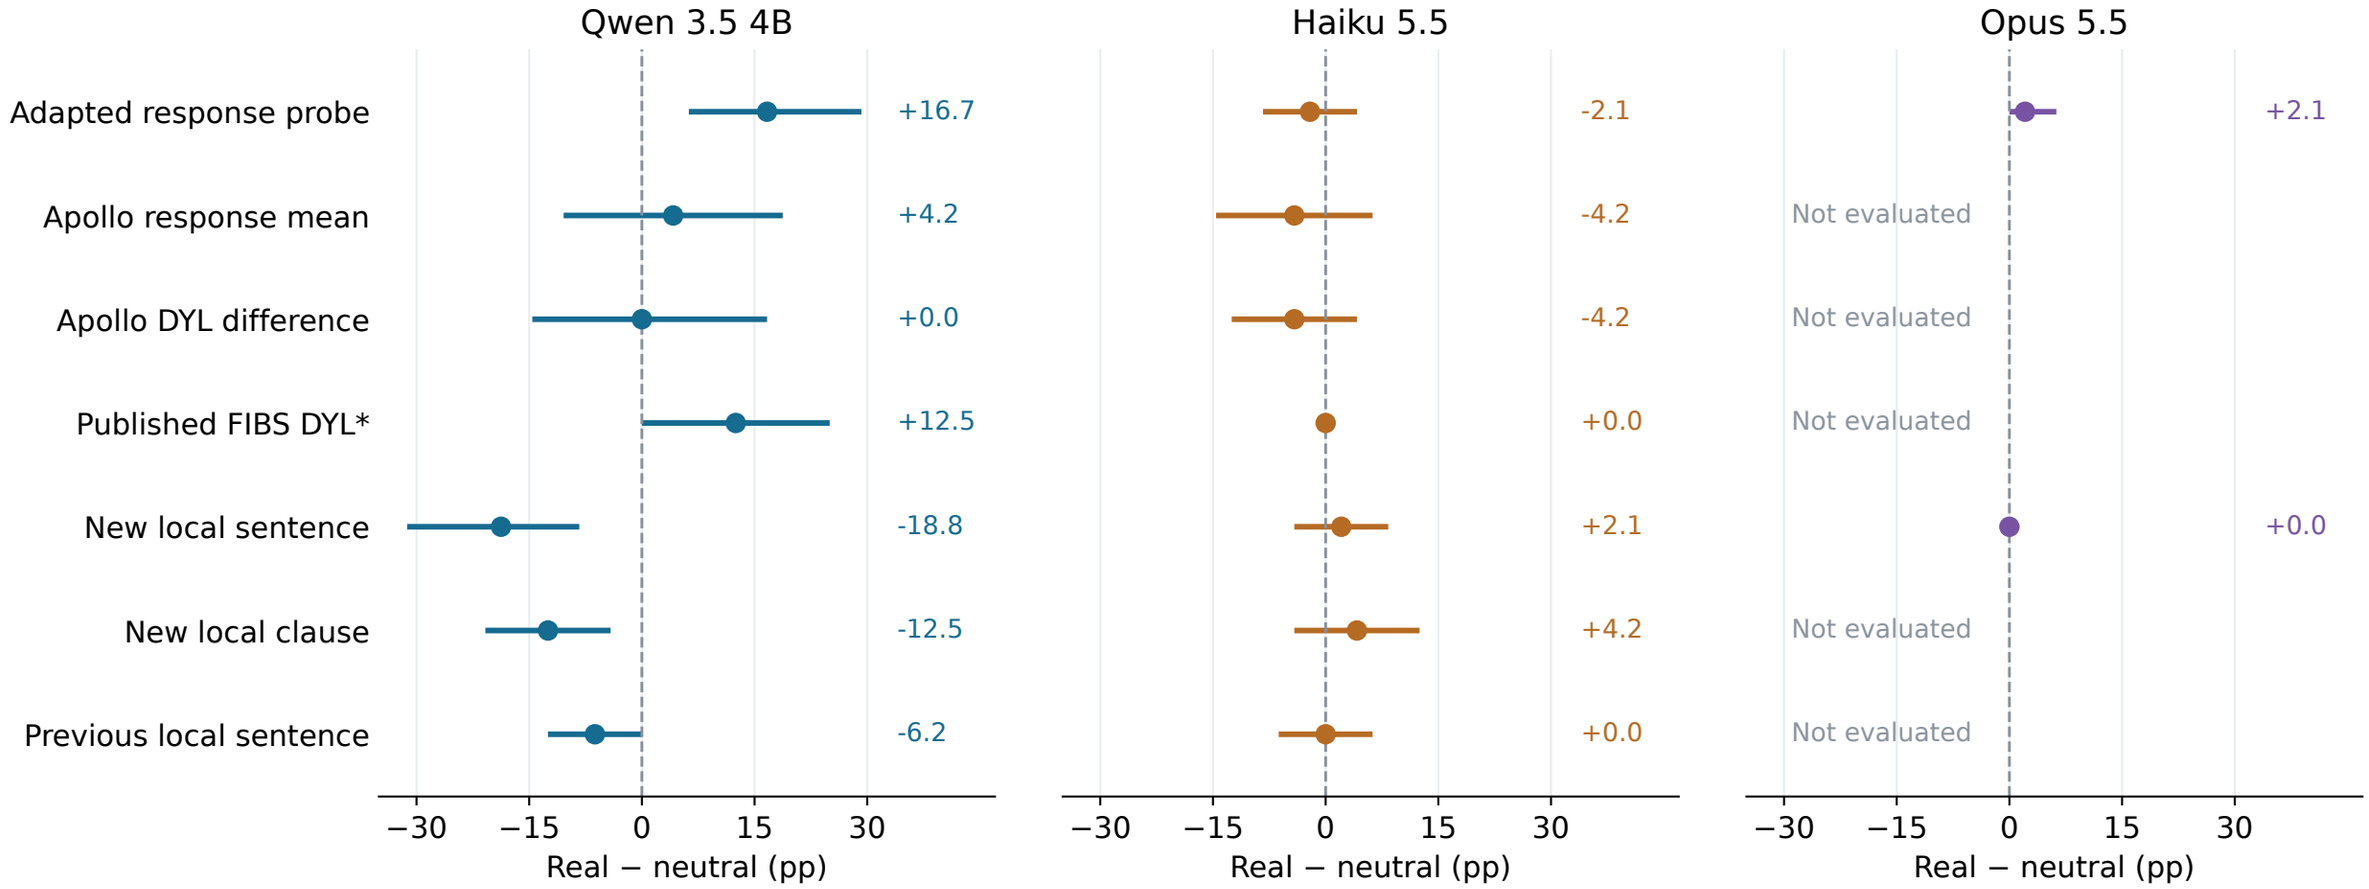
Raw baseline, each probe's neutral control, assisted accuracy, and oracle



48 presentations / 24 story families per cell. Bars: unadjusted 95% family bootstrap intervals.
Oracle is explicitly labelled. A perfect sample score does not establish perfect population accuracy.

All evaluated detector effects

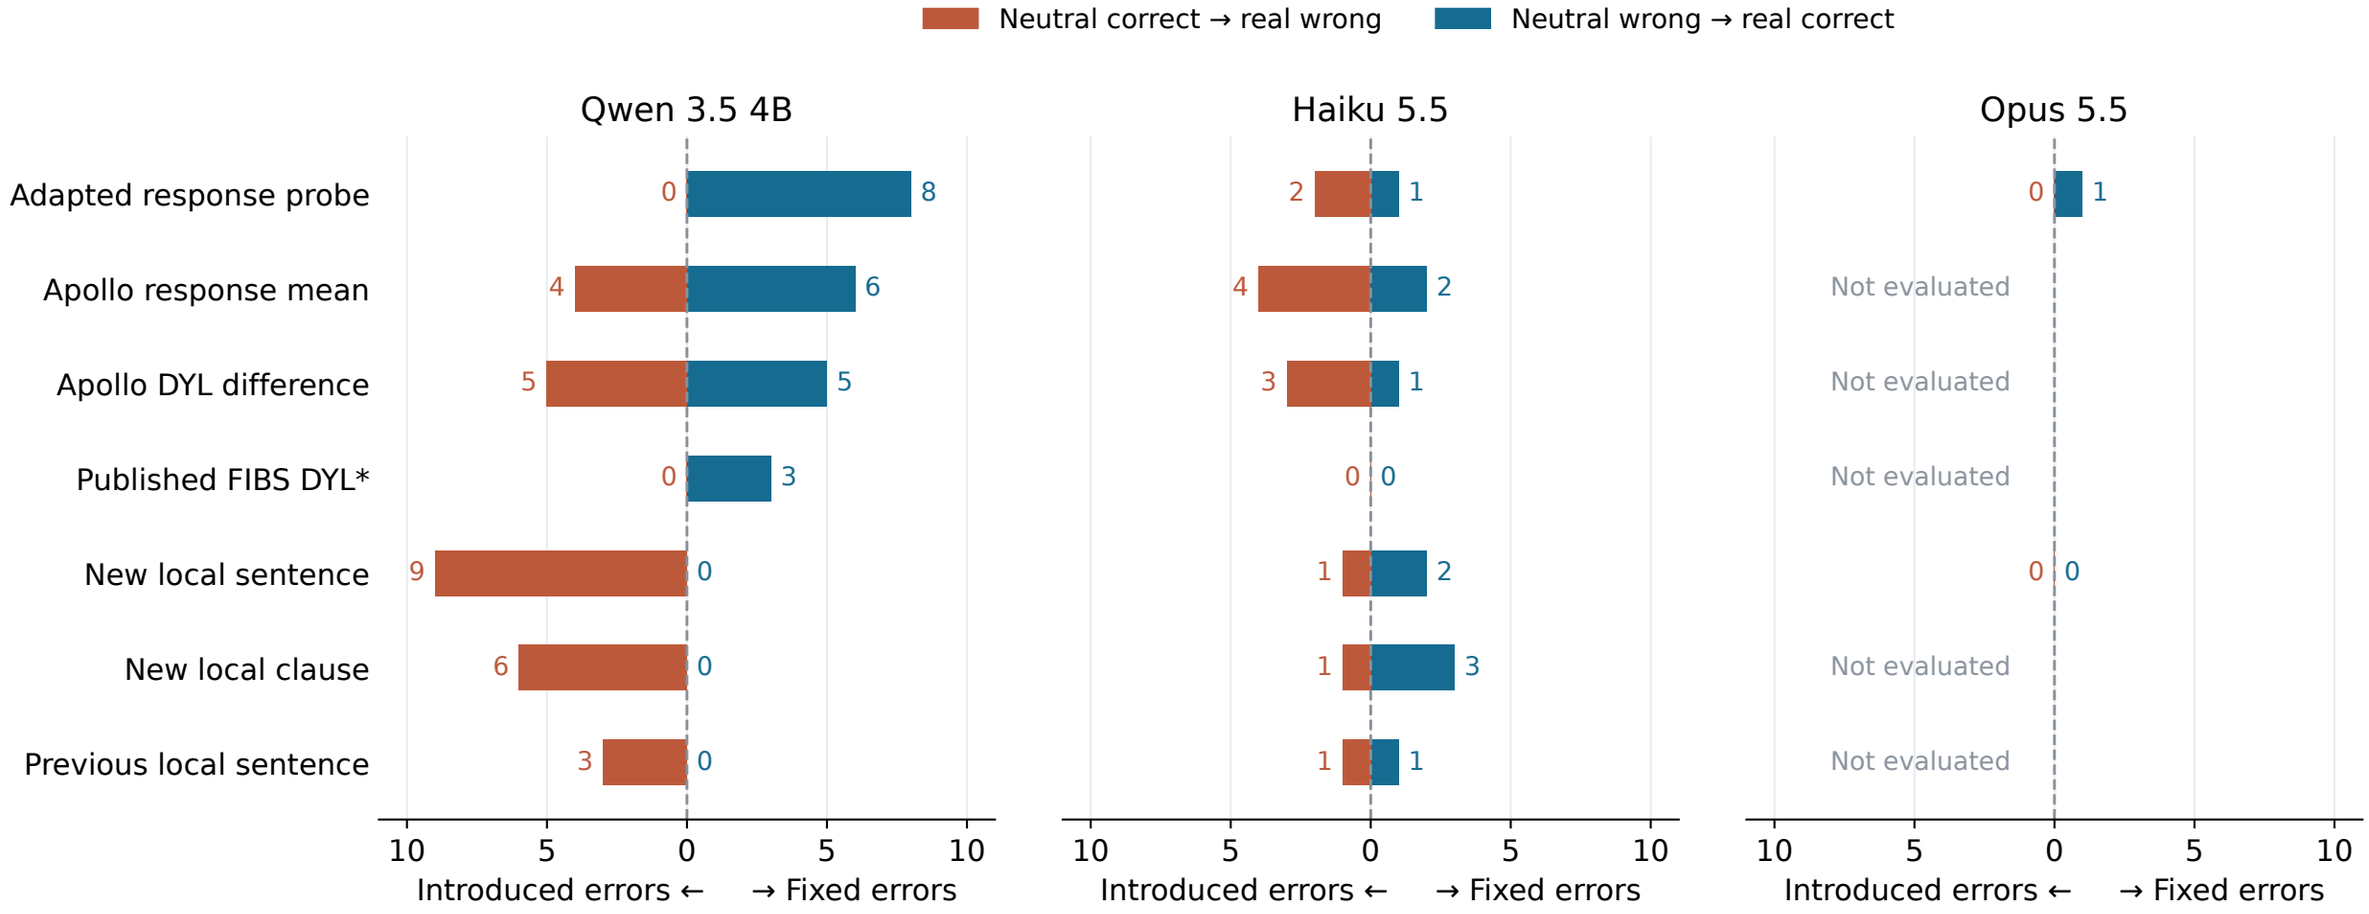
Paired changes in judge accuracy · pooled across debater models where available



Bars: unadjusted 95% source-family bootstrap intervals; 24 independent families. All 16 pooled effects are shown.
*FIBS: Qwen 3.5 debaters only, 24 presentations. Other cells: 48. Opus evaluated only the two prespecified primary signals.

Which answers changed when the probe was shown?

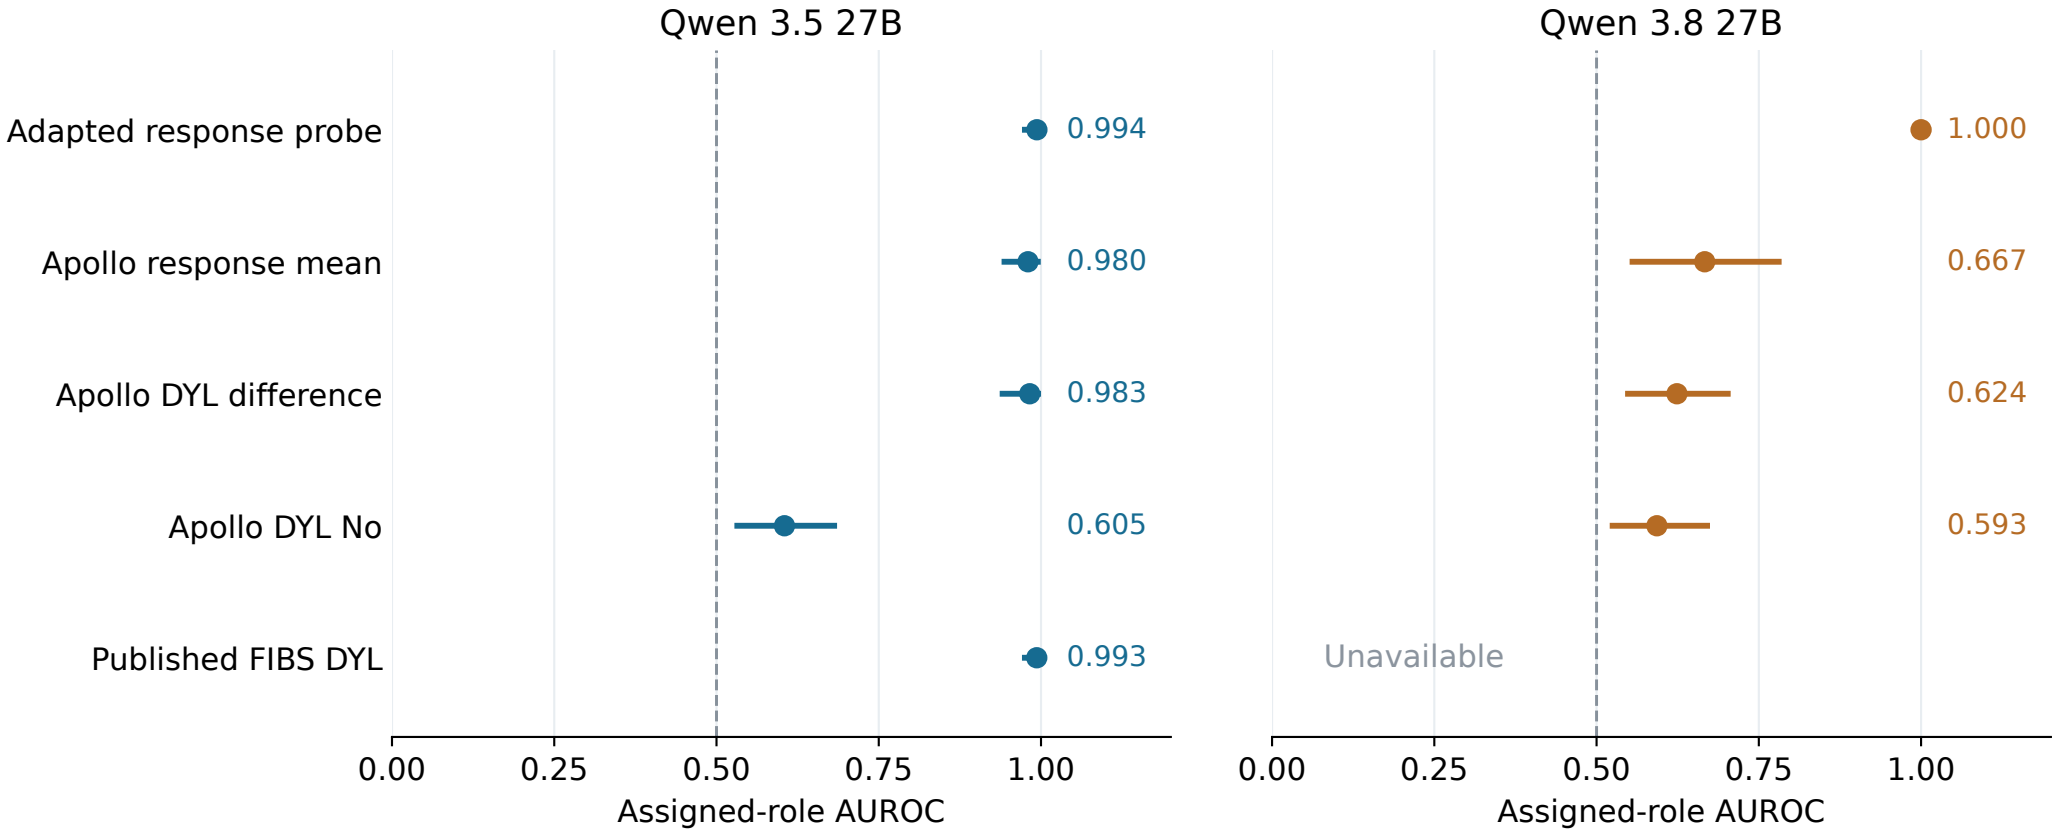
Corrections and new errors relative to the matched neutral display



Counts are paired descriptive outcomes; unchanged answers are omitted. Net uplift = (fixes − harms) / presentations.
48 presentations per cell; *FIBS uses 24 Qwen 3.5 presentations. Original story families, not presentations, are independent.

Whole-response scores distinguish assigned roles

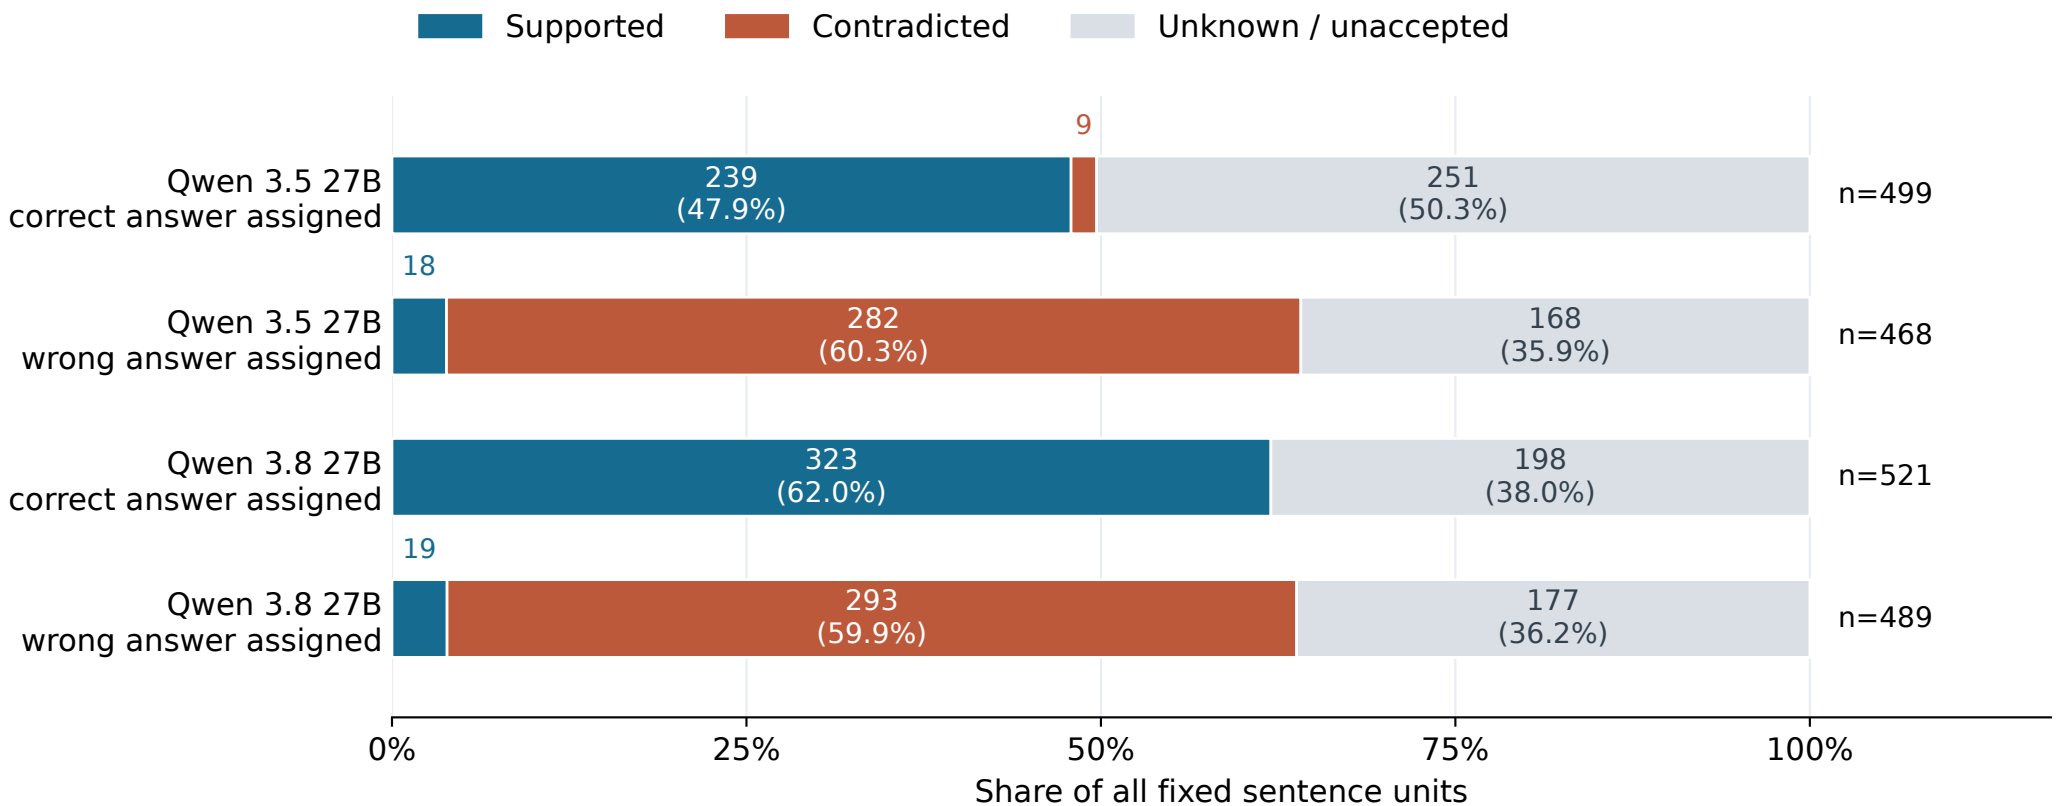
Incorrect-answer assignment versus correct-answer assignment



144 complete turns per debater model, 72 in each role, across 24 original story families. Bars: 95% family bootstrap intervals.
These labels describe assigned roles. They do not establish which individual sentences are false or whether the model intended to deceive.

Assigned role and local truth are different labels

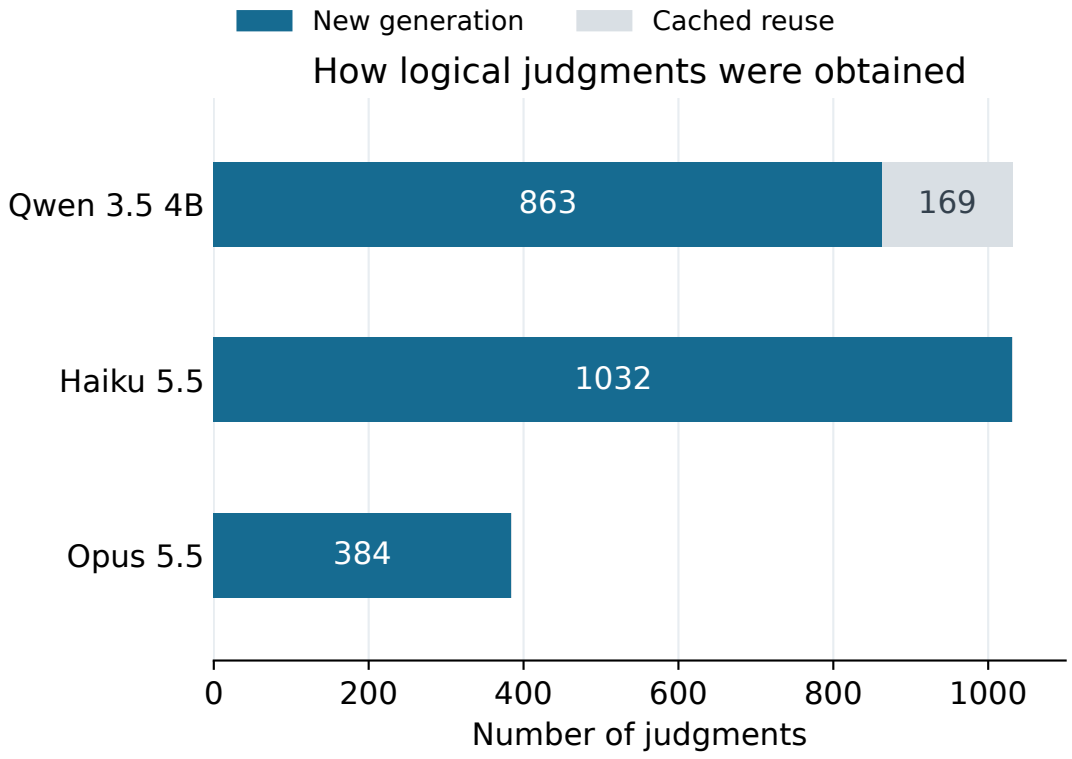
All 1,977 fixed sentence units, including unknowns and annotator disagreement

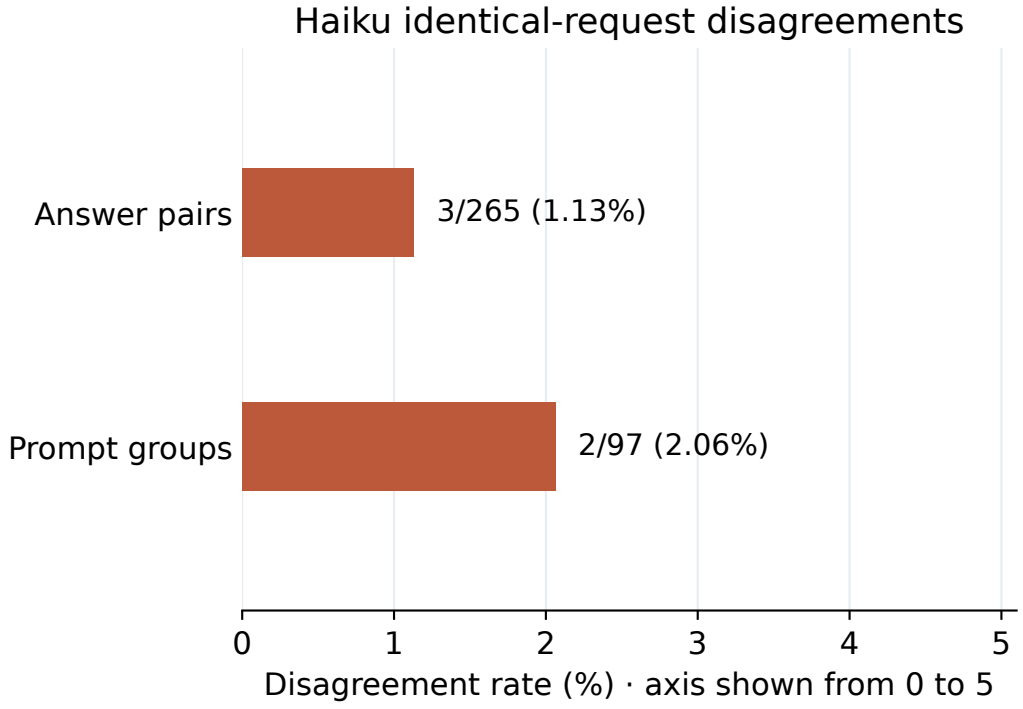


Descriptive unit counts, not independent observations. Each row spans 72 responses from 24 original story families.
Accepted labels require independent model judgments and evidence; they are not human ground truth. Unknown units are not counted as truthful.

The available repeatability check is limited

Cached agreement does not measure generation noise

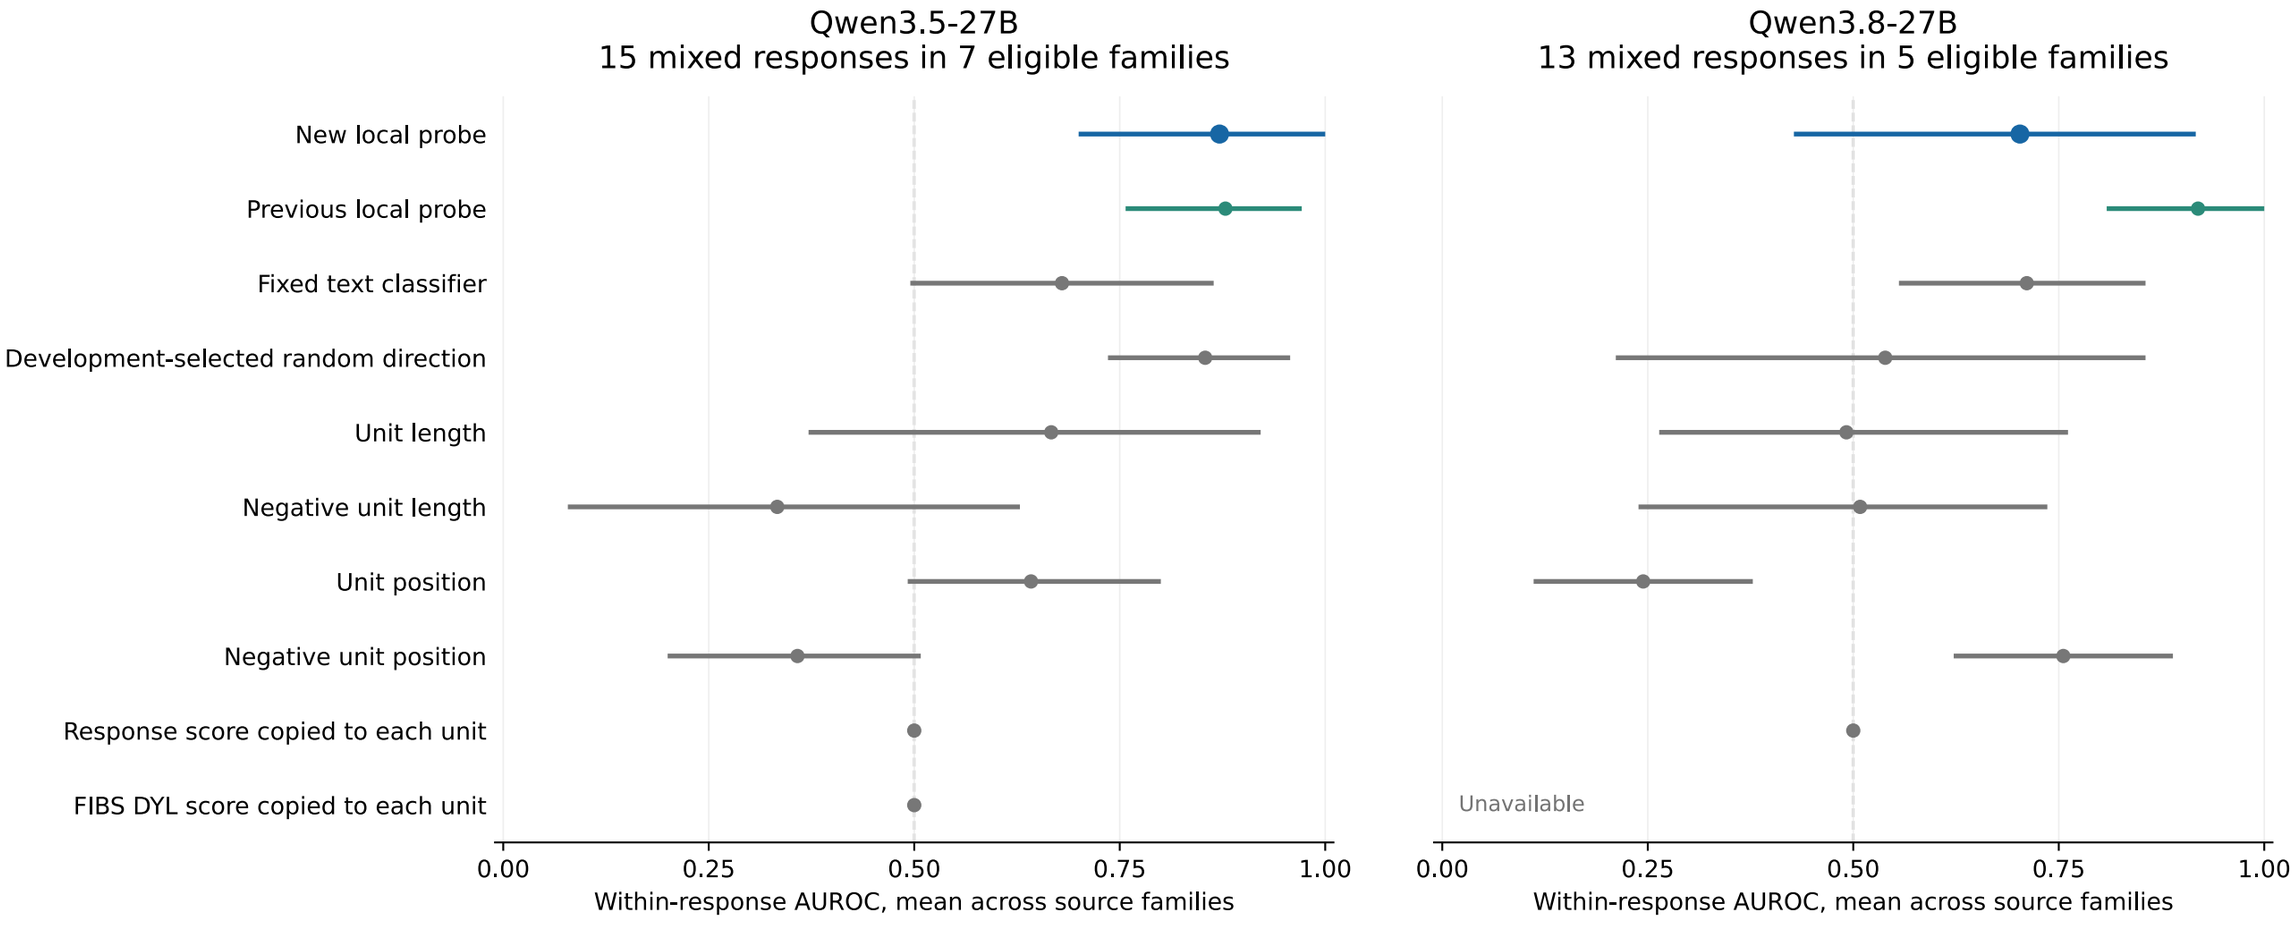




Native Qwen: cached repeats only; fresh-repeat variability unmeasured. Opus: no repeated identical prompts.
Haiku calls span Anthropic/Azure routes without fingerprints. This measures routing/runtime variability, not isolated sampling noise.
Pairs within a prompt group overlap; the two rates are descriptive summaries, not independent trials or confidence intervals.

Primary local evaluation: all 288 turns; accepted mixed-label responses determine within-response ranking.

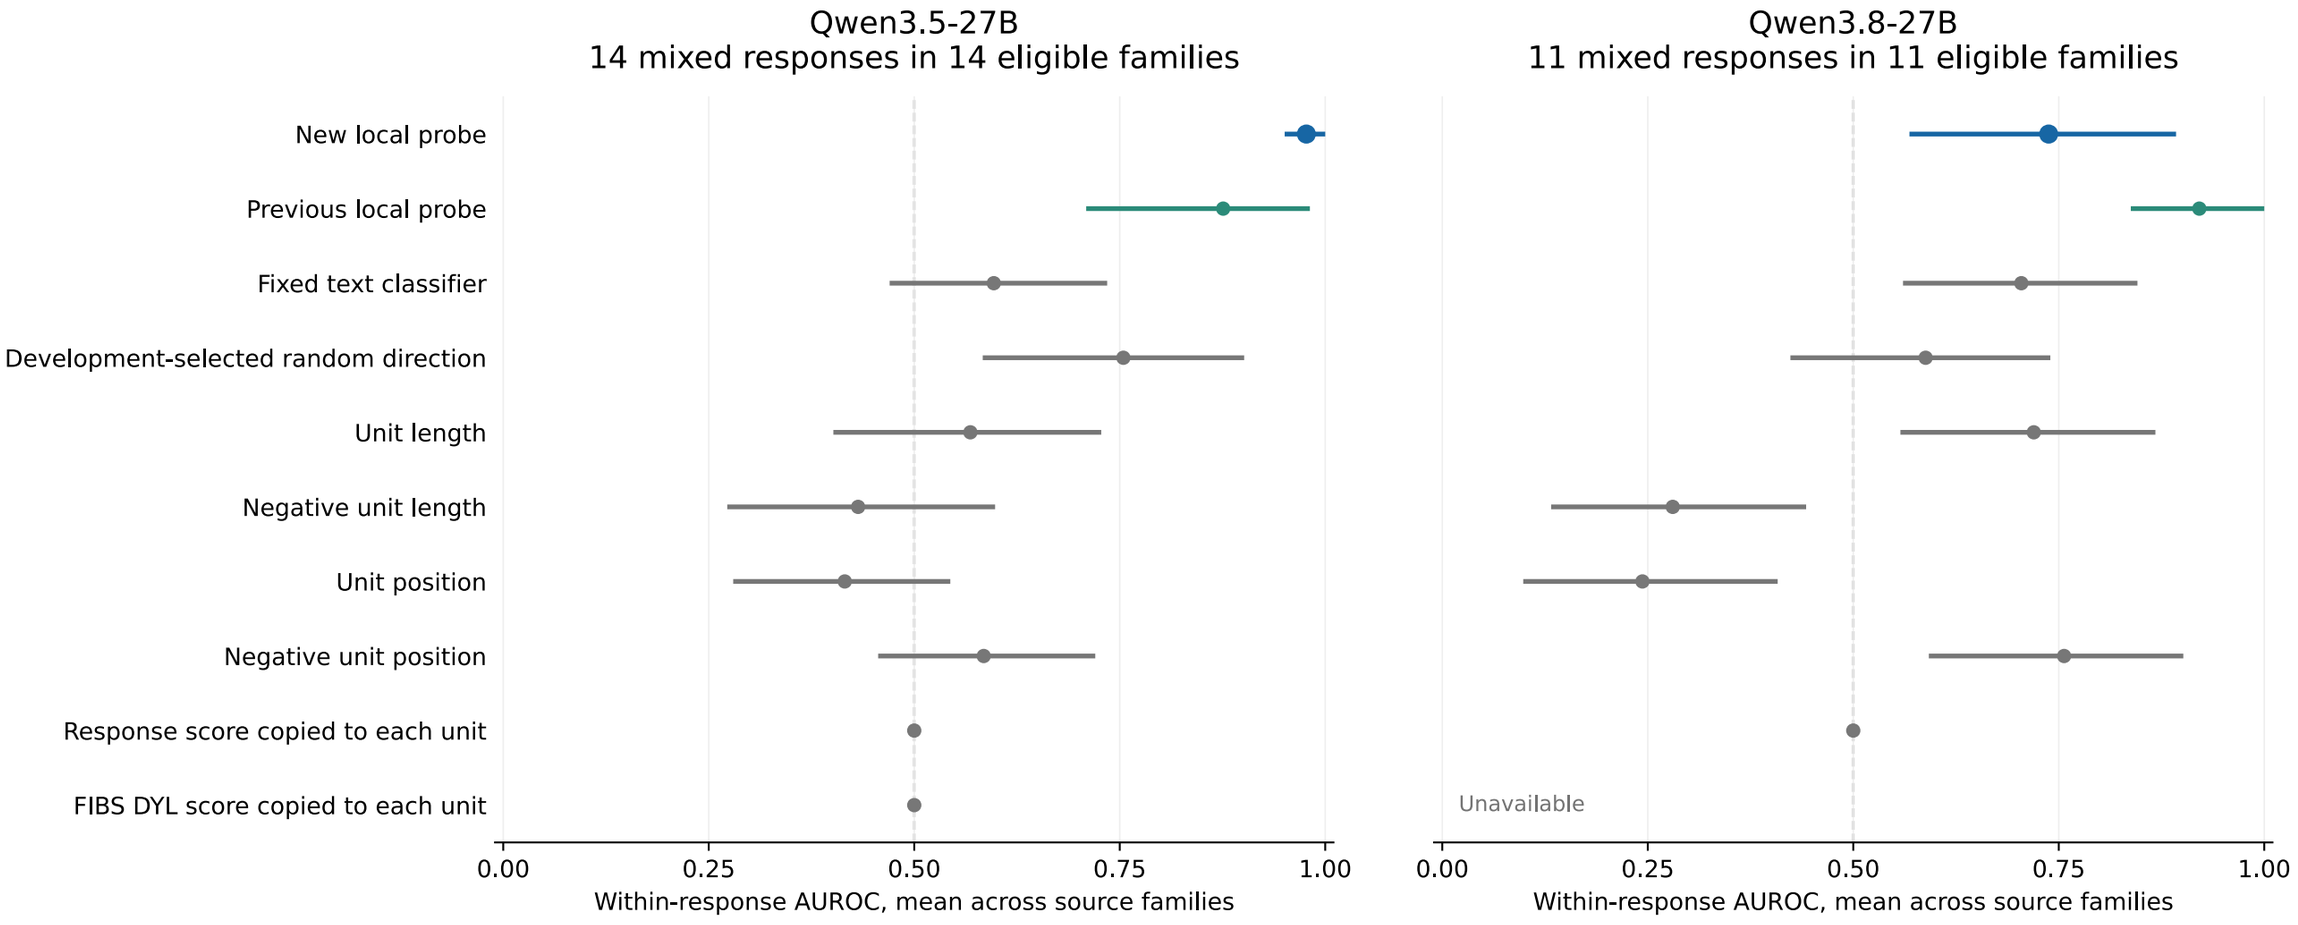
Fresh local probe evaluation — sentence units



Qwen3.5-27B: 548/967 units with accepted truth labels; 144 complete responses; 24 source families.
Qwen3.8-27B: 635/1010 units with accepted truth labels; 144 complete responses; 24 source families.
Intervals resample source families (95%). Only responses containing both accepted supported and contradicted units contribute to local ranking.
Unknown/mixed units stay in coverage counts. Labels are independent model judgments, not human ground truth. All frozen controls are shown.

Secondary diagnostic only: 48 wrong-assigned openings. Distinct from the all-turn clause displays used by judges.

Fresh local probe evaluation — clause units

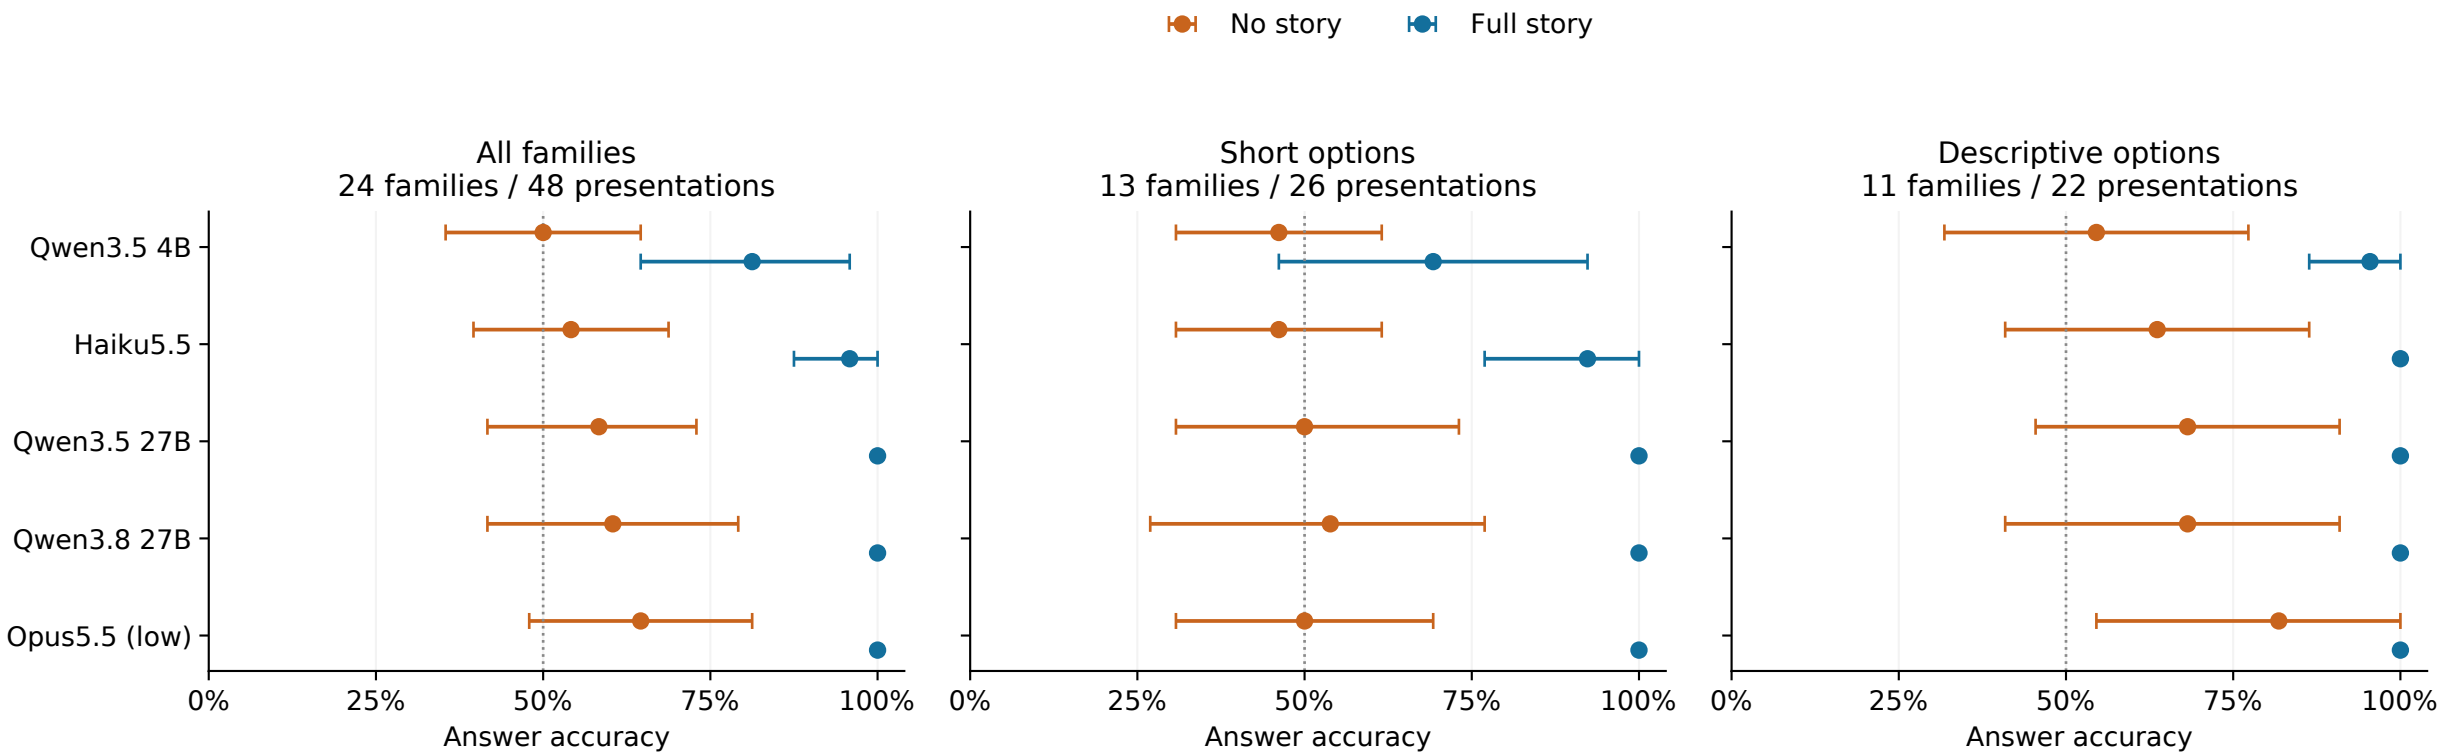


Qwen3.5-27B: 136/391 units with accepted truth labels; 24 complete responses; 24 source families.
Qwen3.8-27B: 142/380 units with accepted truth labels; 24 complete responses; 24 source families.
Intervals resample source families (95%). Only responses containing both accepted supported and contradicted units contribute to local ranking.
Unknown/mixed units stay in coverage counts. Labels are independent model judgments, not human ground truth. All frozen controls are shown.

Fresh corrected source controls: question-only versus full-story answers, with short-option subgroup reported separately.

Source-reading controls after source metadata cleanup

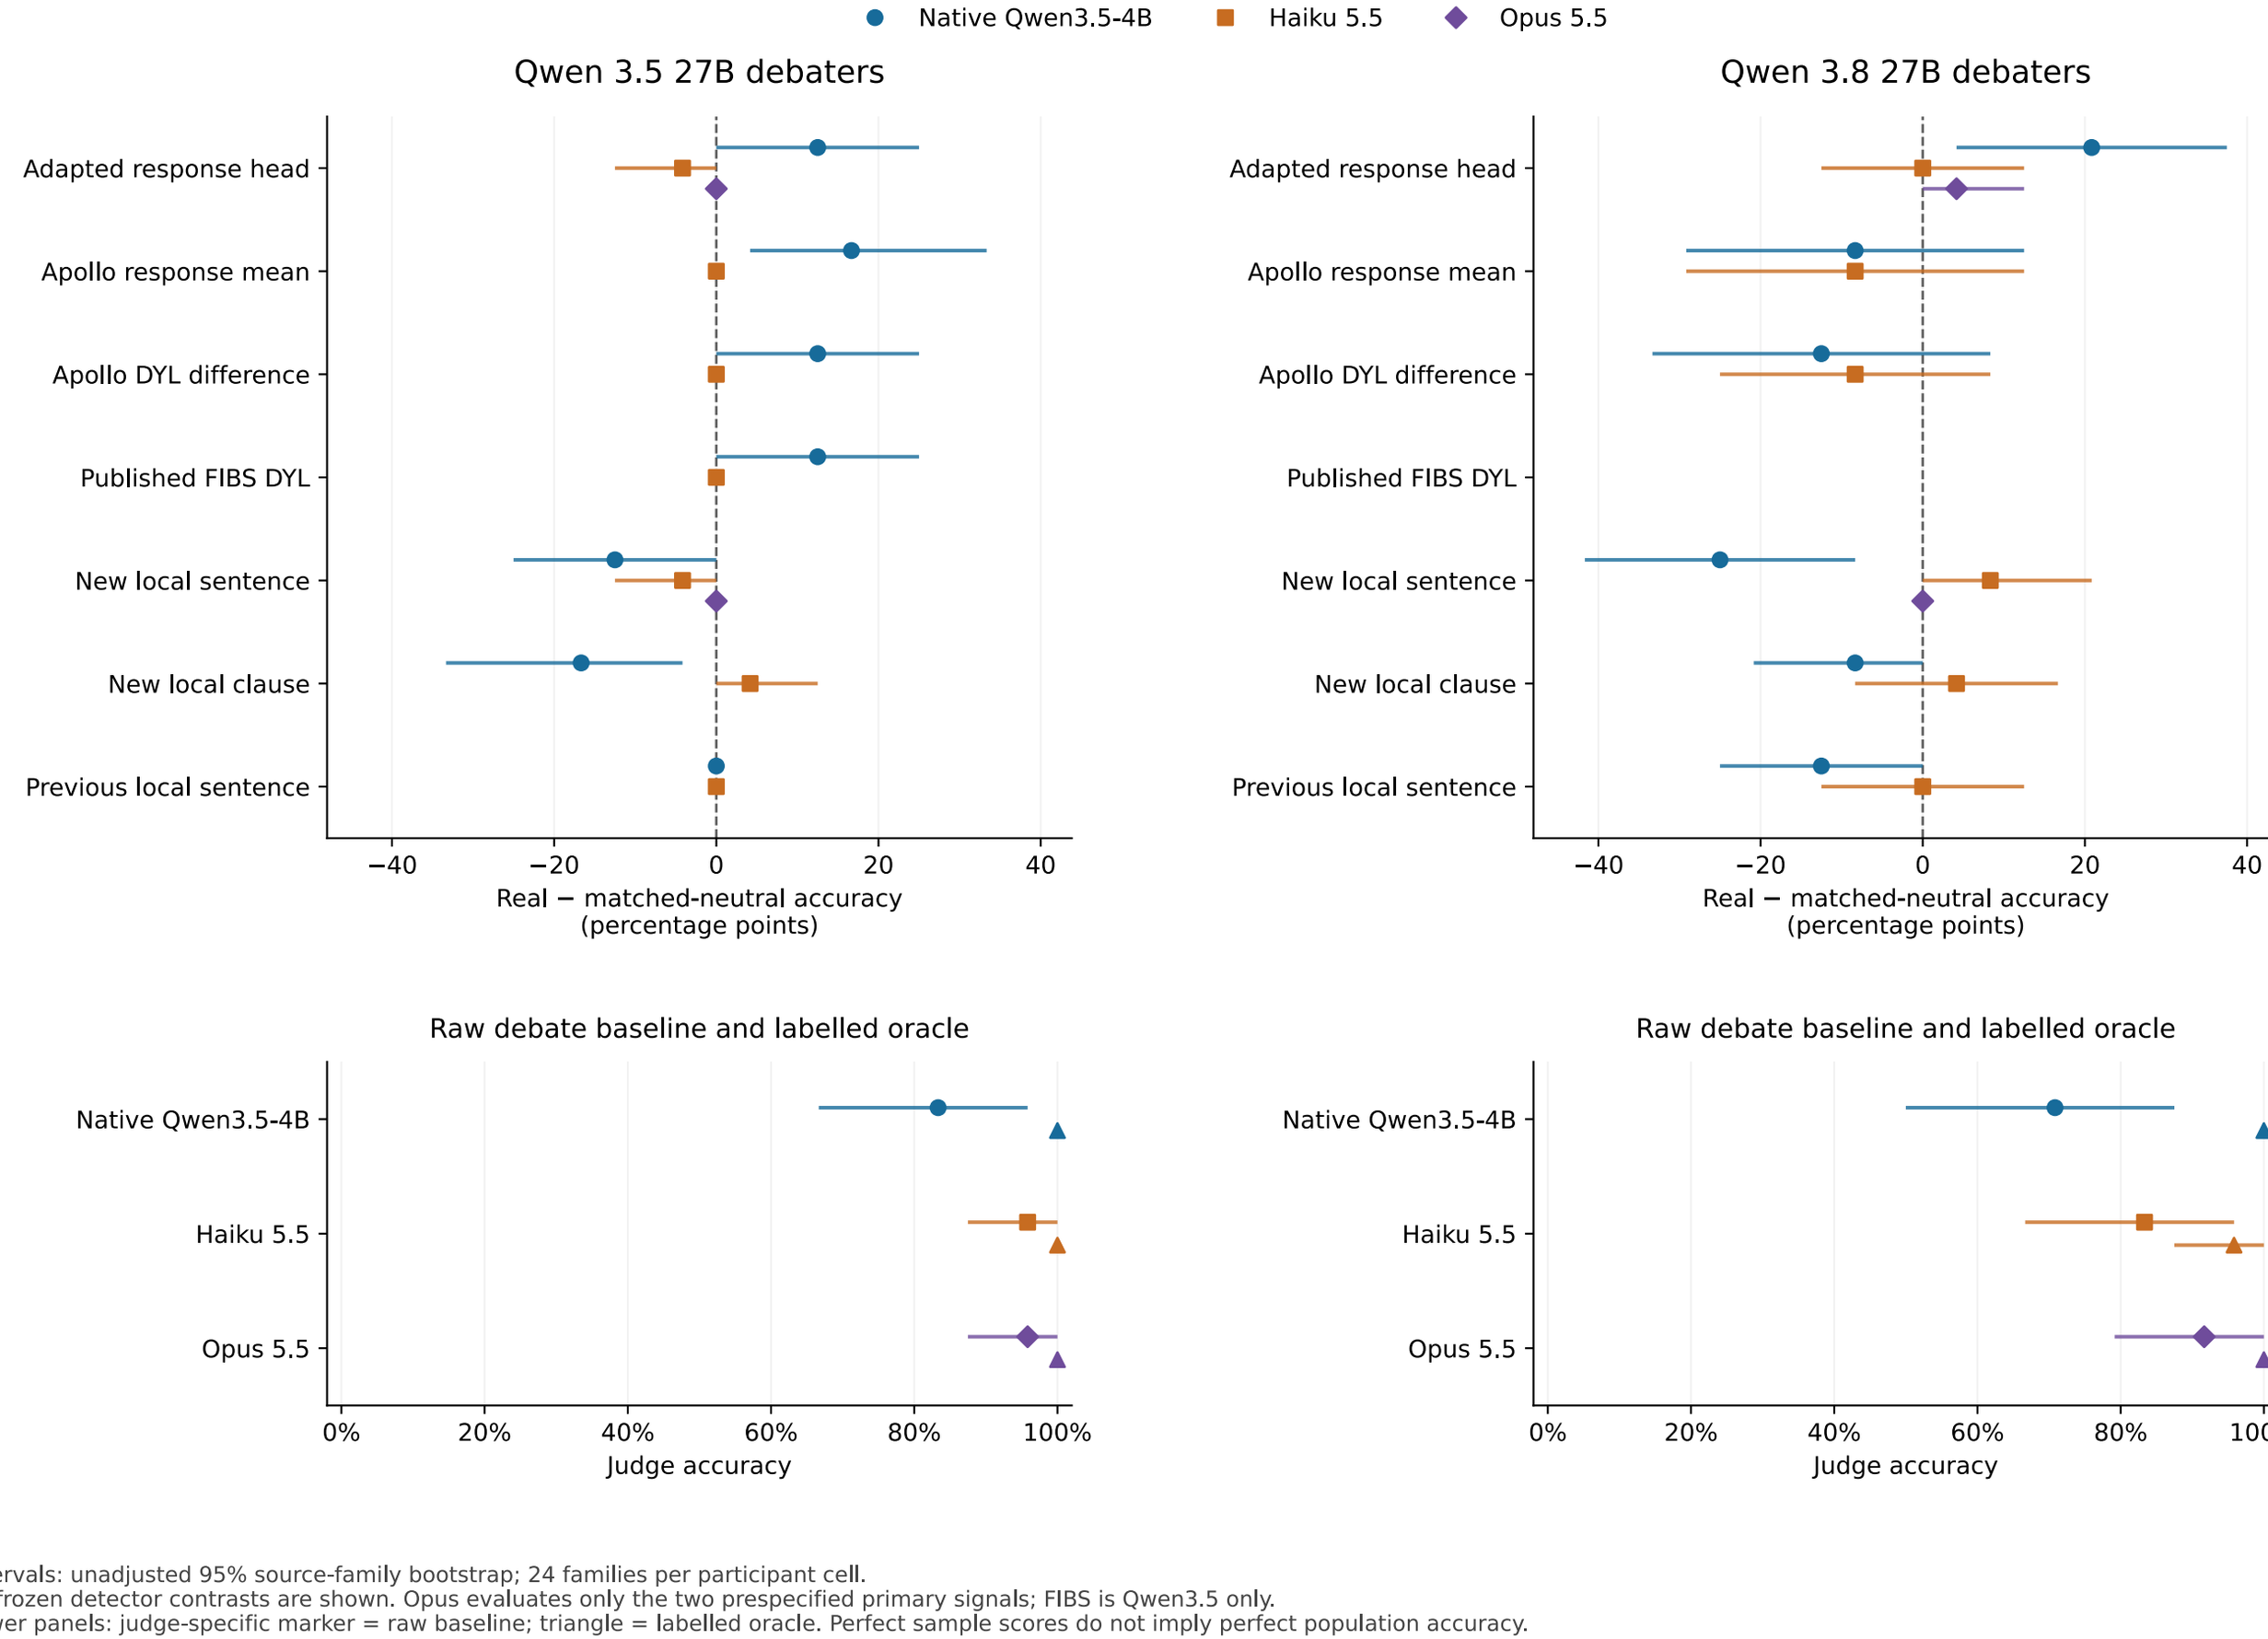
All 24 held-out source families retained; both answer orders



Points: observed accuracy. Bars: 95% bootstrap intervals over source families (10,000 draws).
Only source 114's full-story controls were rerun after removing an appended answer-key footer; original records remain intact.
Source controls do not measure debate judging or probe uplift. Perfect sample accuracy does not imply population certainty.

Detailed judge comparisons: every frozen participant-specific effect, plus raw no-signal baselines and labelled oracle.

Probity: frozen probe assistance comparison



Intervals: unadjusted 95% source-family bootstrap; 24 families per participant cell.
All frozen detector contrasts are shown. Opus evaluates only the two prespecified primary signals; FIBS is Qwen3.5 only.
Lower panels: judge-specific marker = raw baseline; triangle = labelled oracle. Perfect sample scores do not imply perfect population accuracy.