Judge accuracy across the completed 2×2 grid
166 questions · 116 source families · one trajectory seed

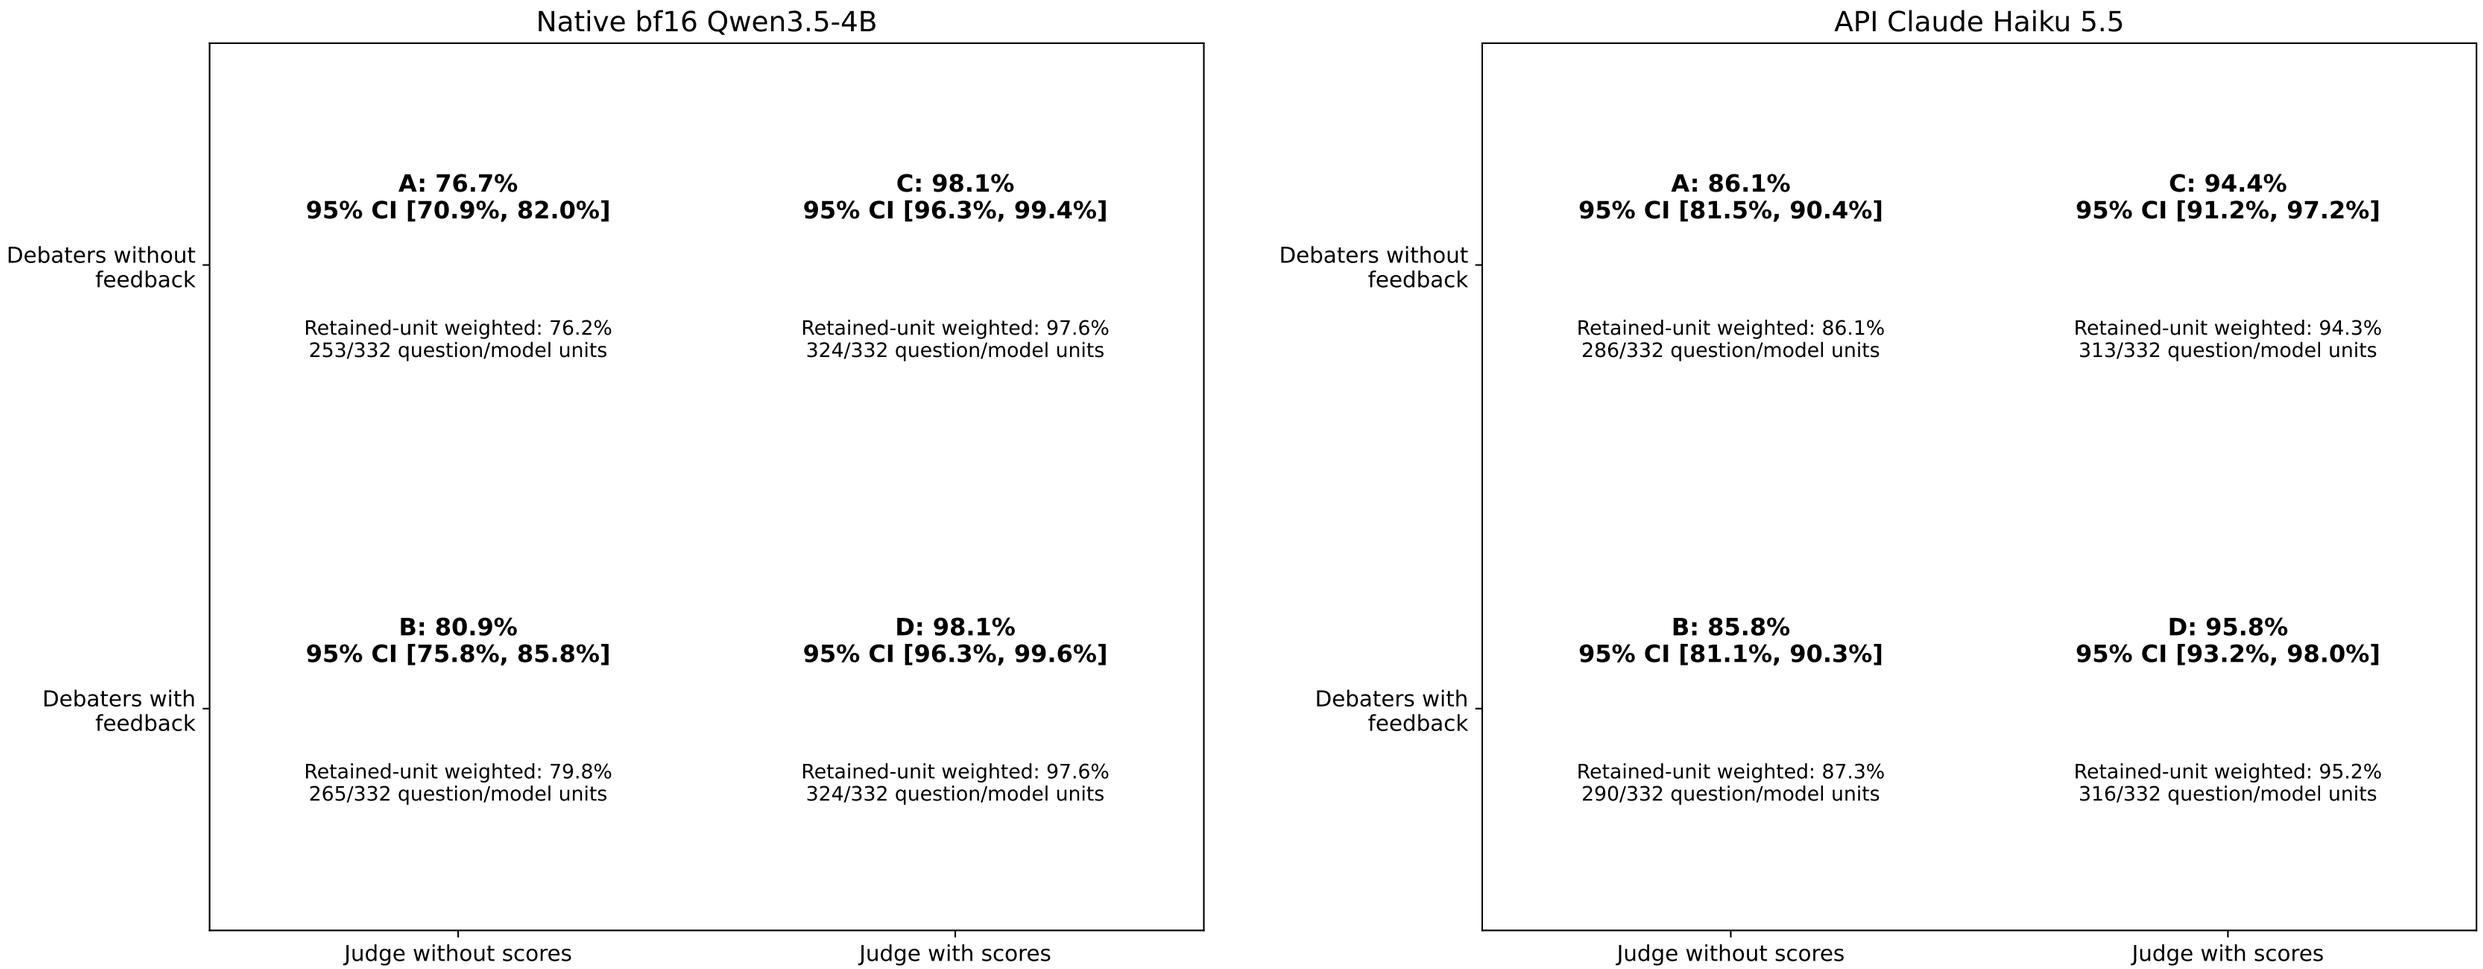

A: no feedback / no packet · B: feedback-conditioned / no packet · C: exact A transcript / packet · D: exact B transcript / packet.

Packet: frozen adapted response-mean logistic head, whole structured-response mean pooling; Qwen3.5 layer 43 / Qwen3.8 layer 35; prior-turn percentiles, speaker means, optional recommendation.

Family-macro headline; retained-unit rate shown separately. Intervals are unadjusted 95% percentile family bootstrap, 10,000 resamples; no repeat-generation uncertainty.

Full corpus — descriptive paired comparisons

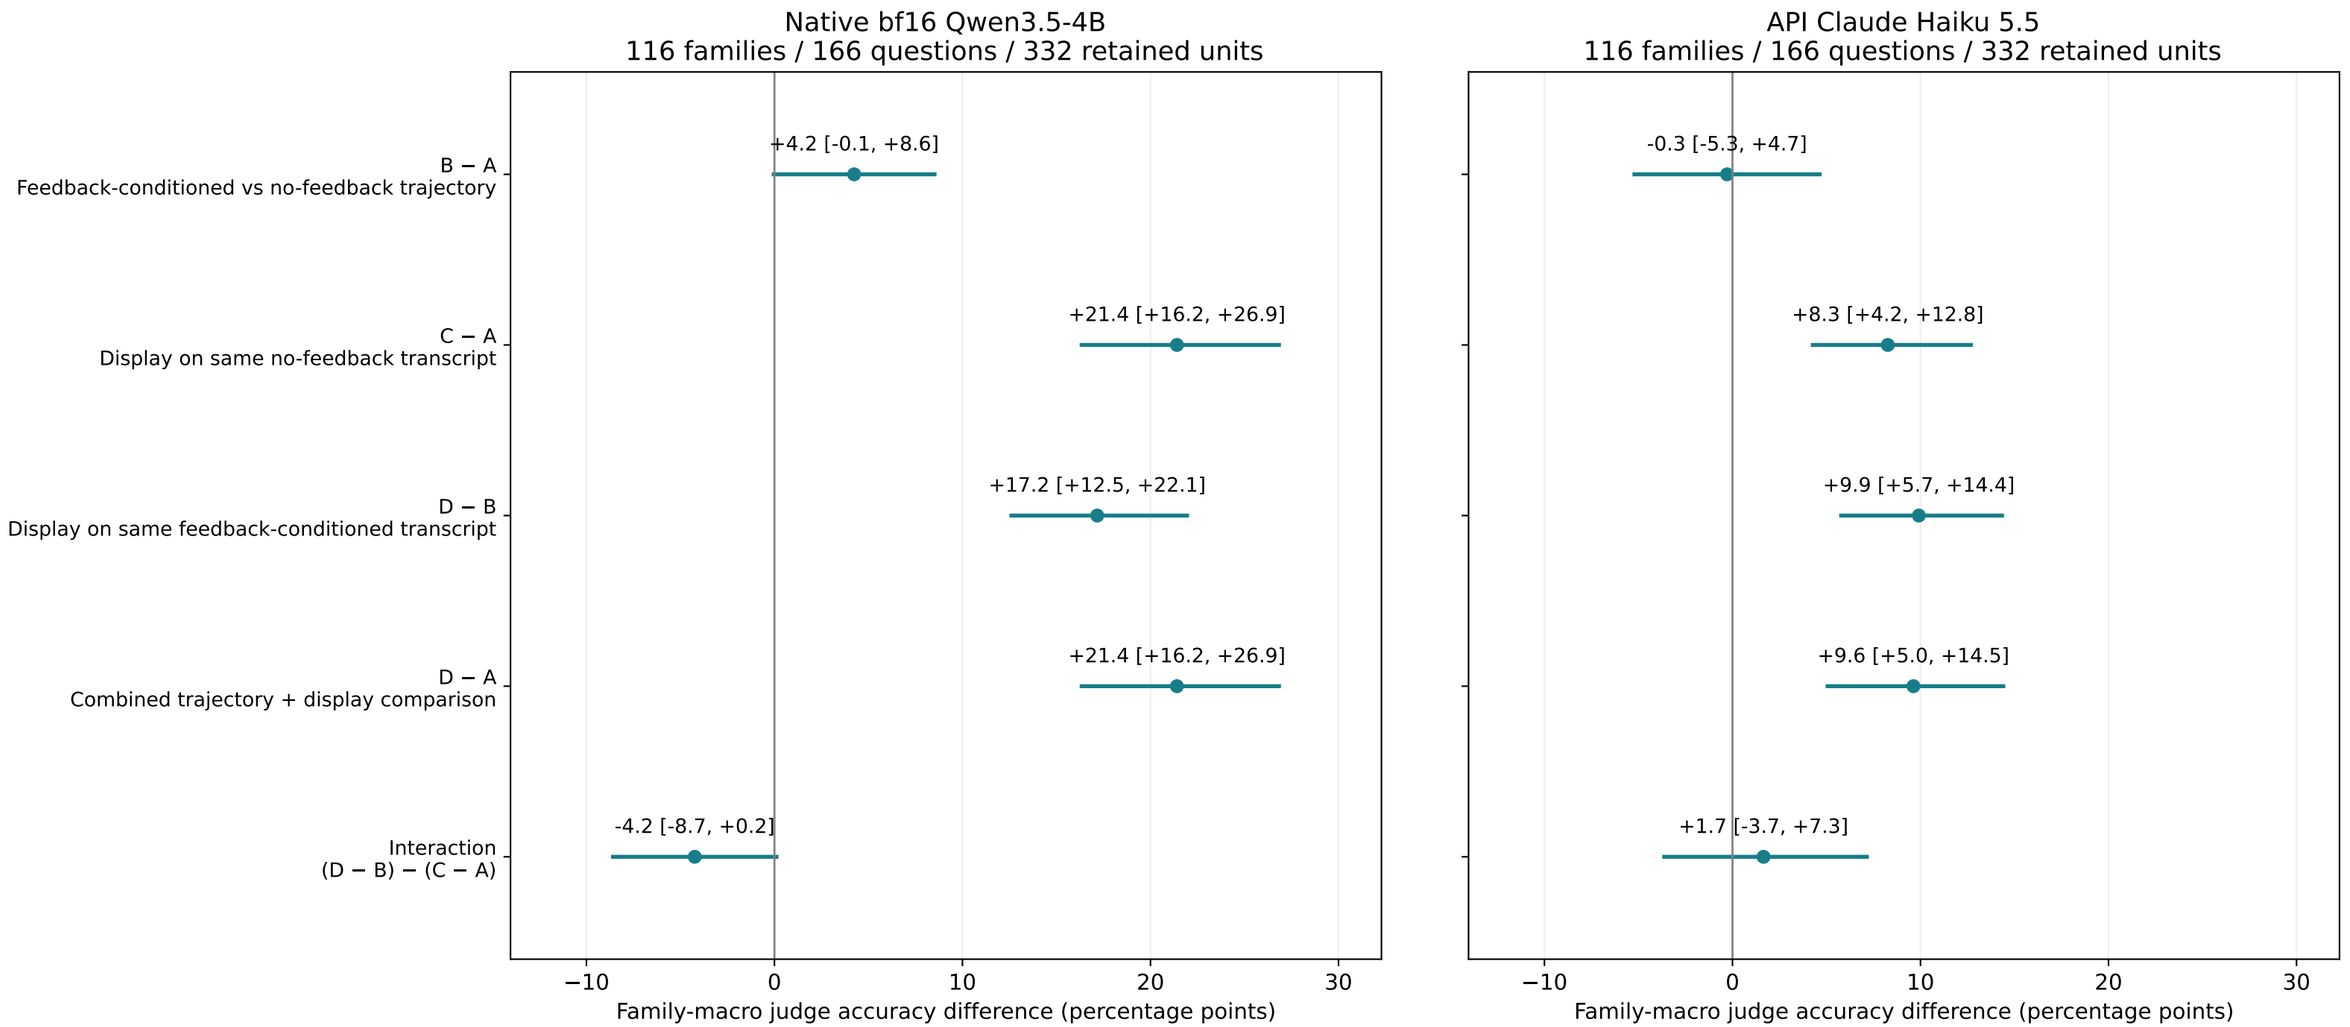

A: no feedback / no packet · B: feedback-conditioned / no packet · C: exact A transcript / packet · D: exact B transcript / packet.
Packet: frozen adapted response-mean logistic head, whole structured-response mean pooling; Qwen3.5 layer 43 / Qwen3.8 layer 35; prior-turn percentiles, speaker means, optional recommendation.
Equal-weight source-family means. Intervals: unadjusted 95% percentile family bootstrap, 10,000 resamples.
They describe family variation; one debate seed gives no repeated-generation uncertainty. B −A/D−A change trajectories; C−A/D−B hold public text fixed.

Families untouched before expansion — descriptive sensitivity

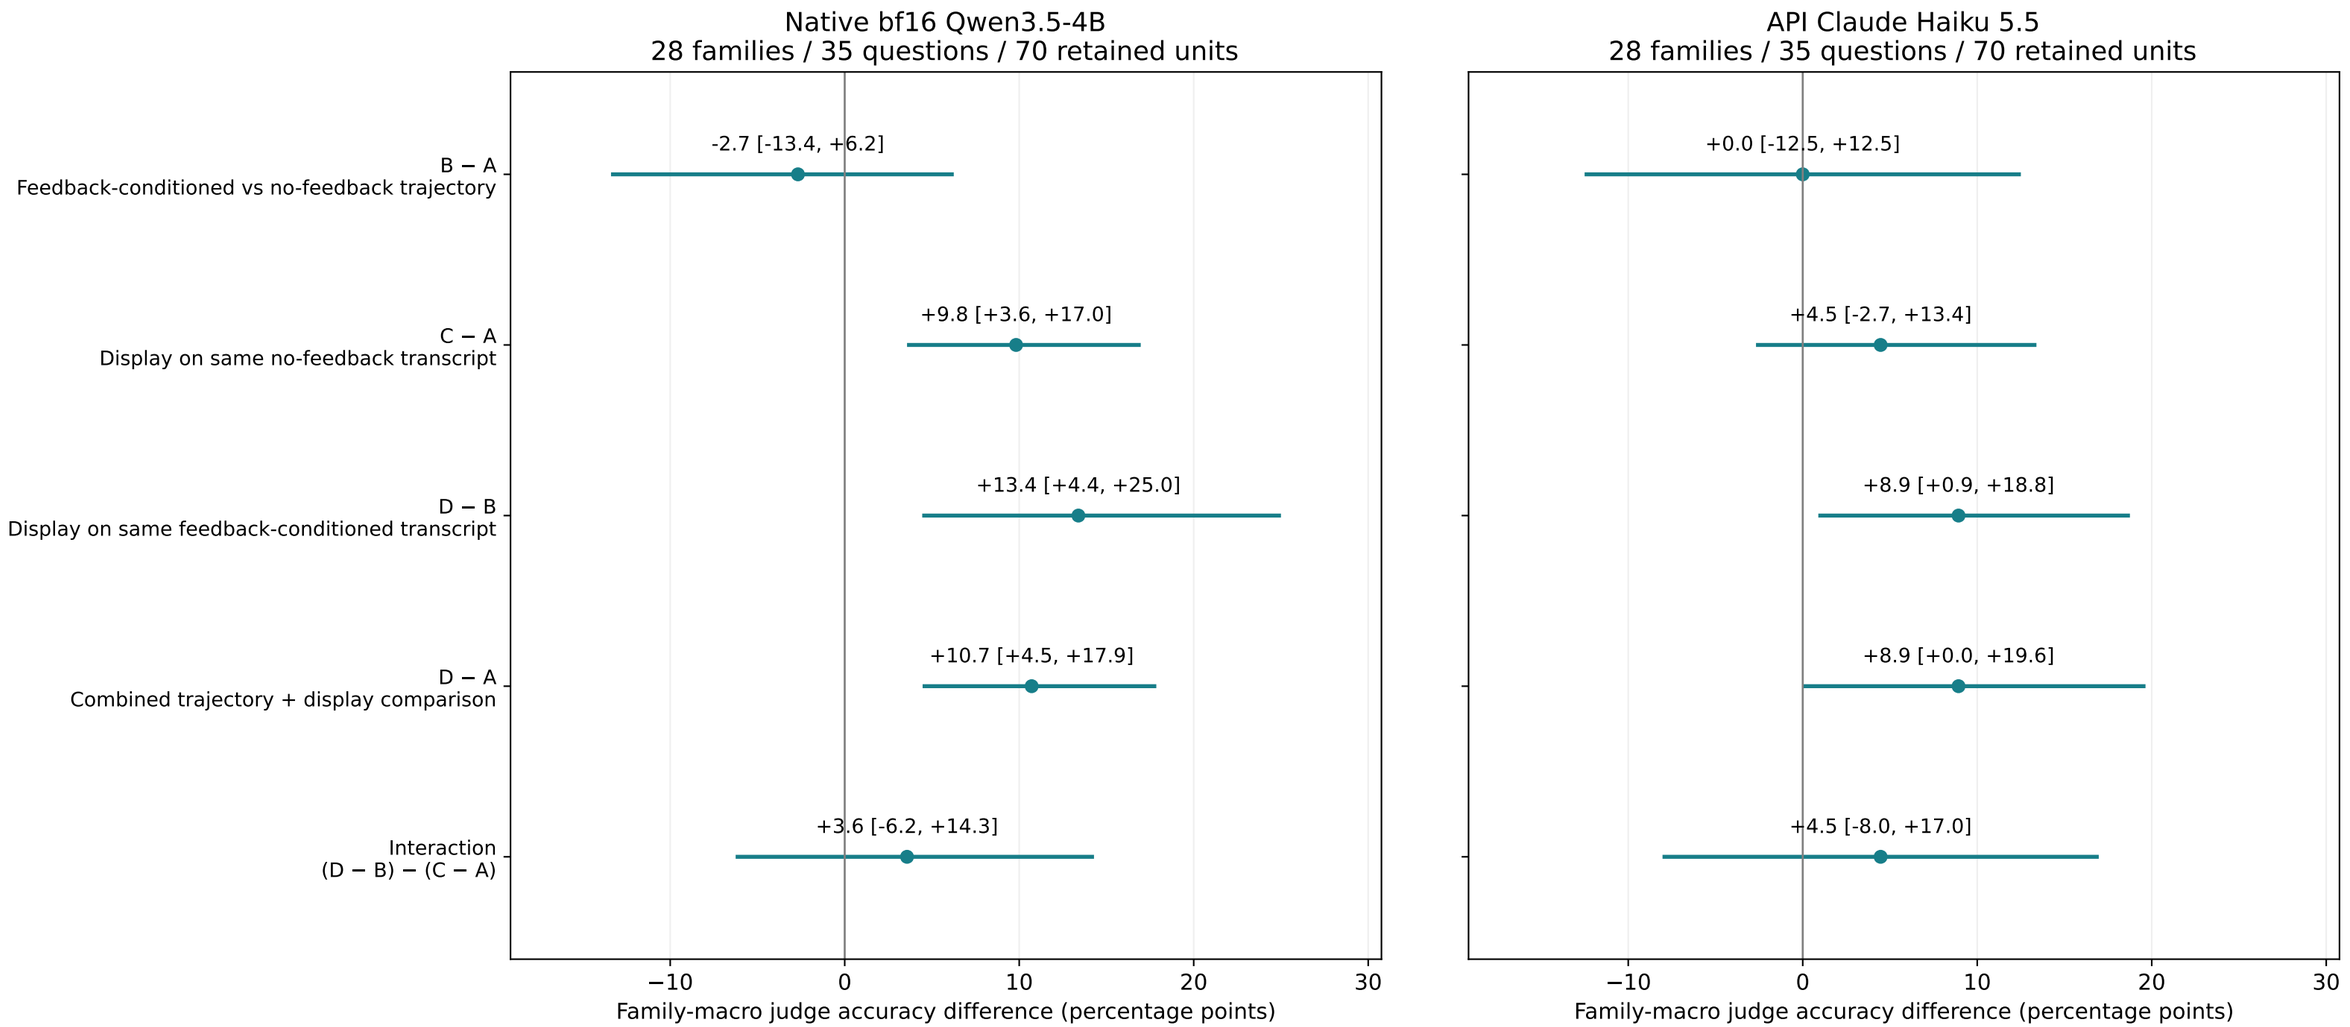

A: no feedback / no packet · B: feedback-conditioned / no packet · C: exact A transcript / packet · D: exact B transcript / packet.
Packet: frozen adapted response-mean logistic head, whole structured-response mean pooling; Qwen3.5 layer 43 / Qwen3.8 layer 35; prior-turn percentiles, speaker means, optional recommendation.
Equal-weight source-family means. Intervals: unadjusted 95% percentile family bootstrap, 10,000 resamples.
They describe family variation; one debate seed gives no repeated-generation uncertainty. B −A/D−A change trajectories; C−A/D−B hold public text fixed.

Sensitivity: source-reviewer-supported questions

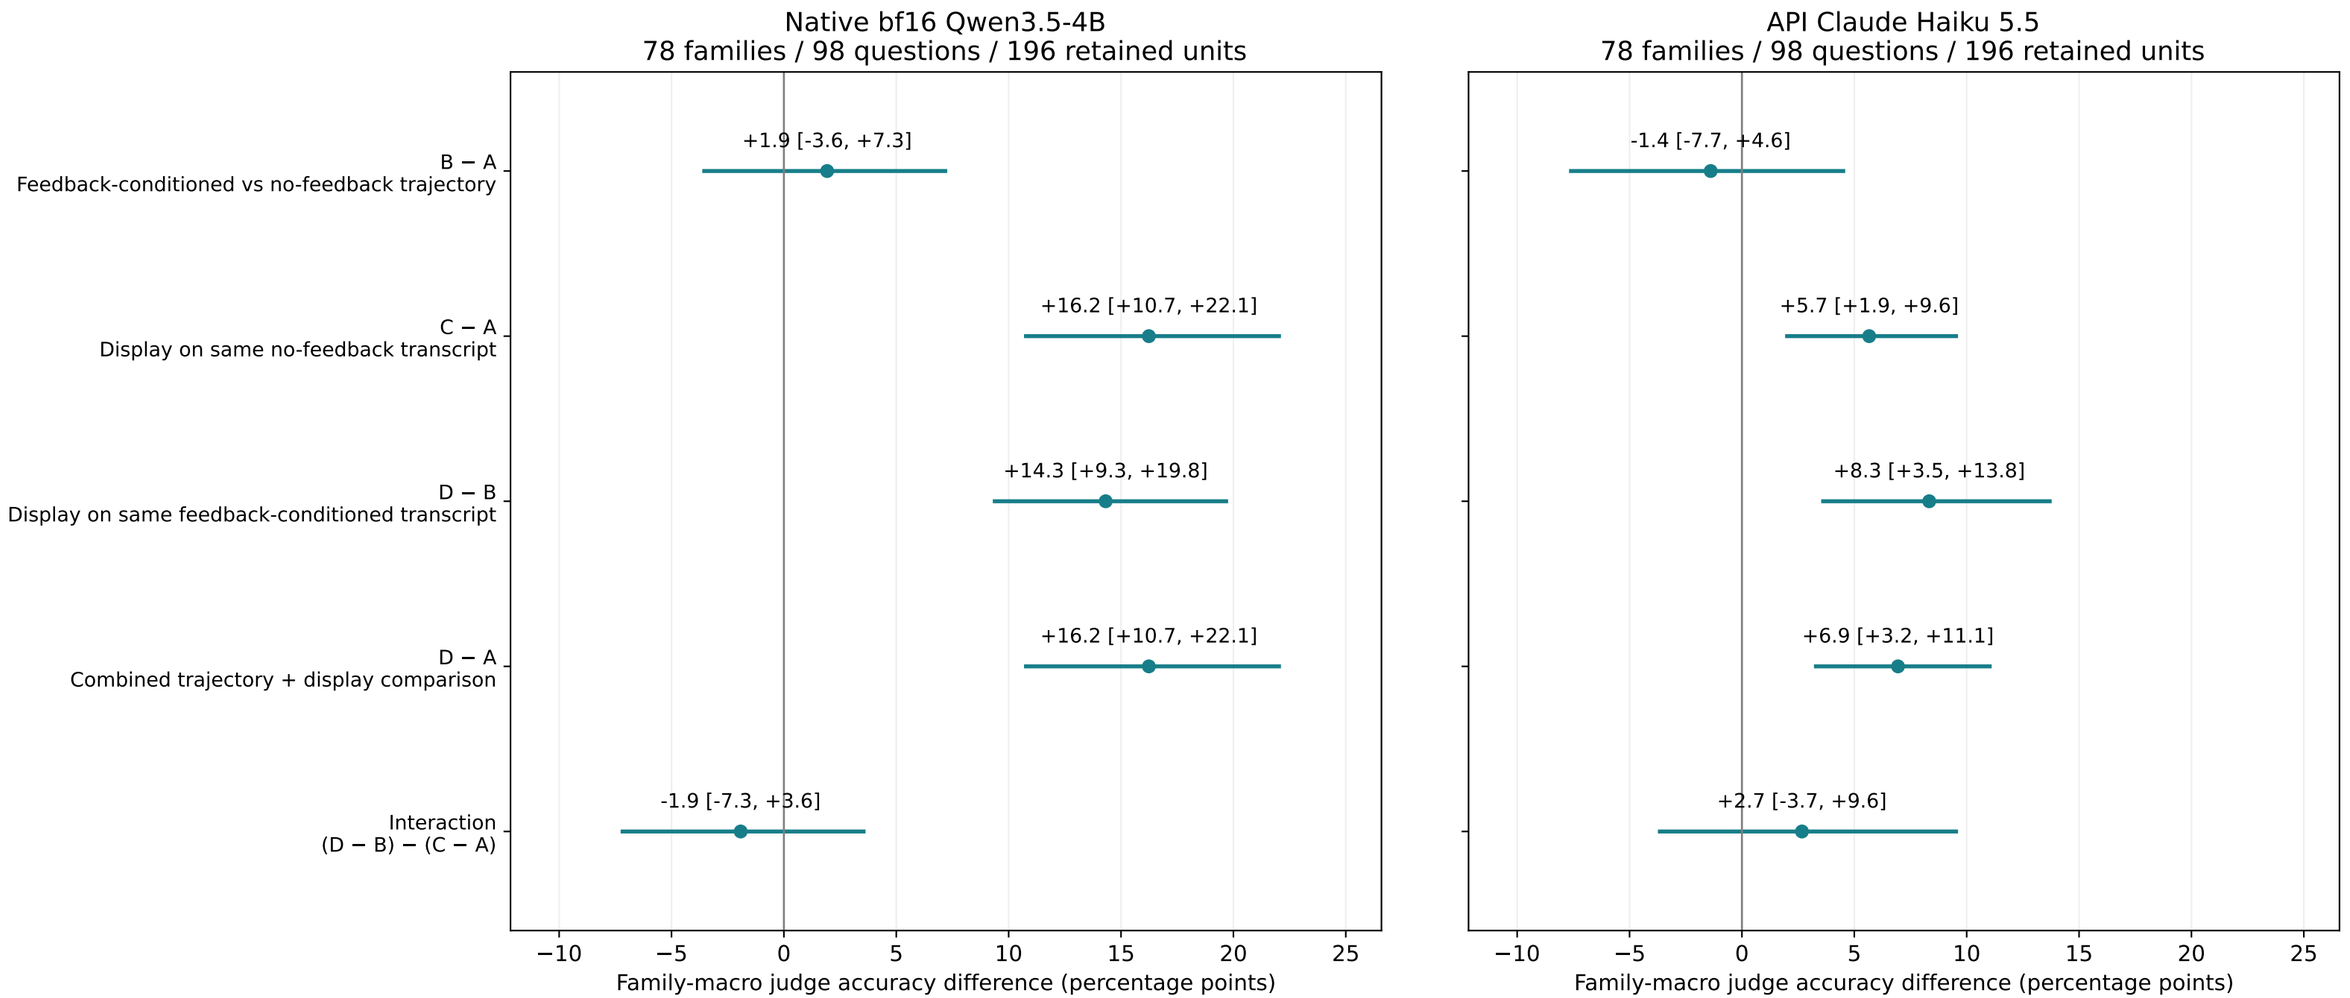

A: no feedback / no packet · B: feedback-conditioned / no packet · C: exact A transcript / packet · D: exact B transcript / packet.
Packet: frozen adapted response-mean logistic head, whole structured-response mean pooling; Qwen3.5 layer 43 / Qwen3.8 layer 35; prior-turn percentiles, speaker means, optional recommendation.
Equal-weight source-family means. Intervals: unadjusted 95% percentile family bootstrap, 10,000 resamples.
They describe family variation; one debate seed gives no repeated-generation uncertainty. B −A/D−A change trajectories; C−A/D−B hold public text fixed.
Selection note: source support is a fallible reviewer diagnostic, not certified gold or local factual truth.

Joint sensitivity: source support AND model-specific eligibility

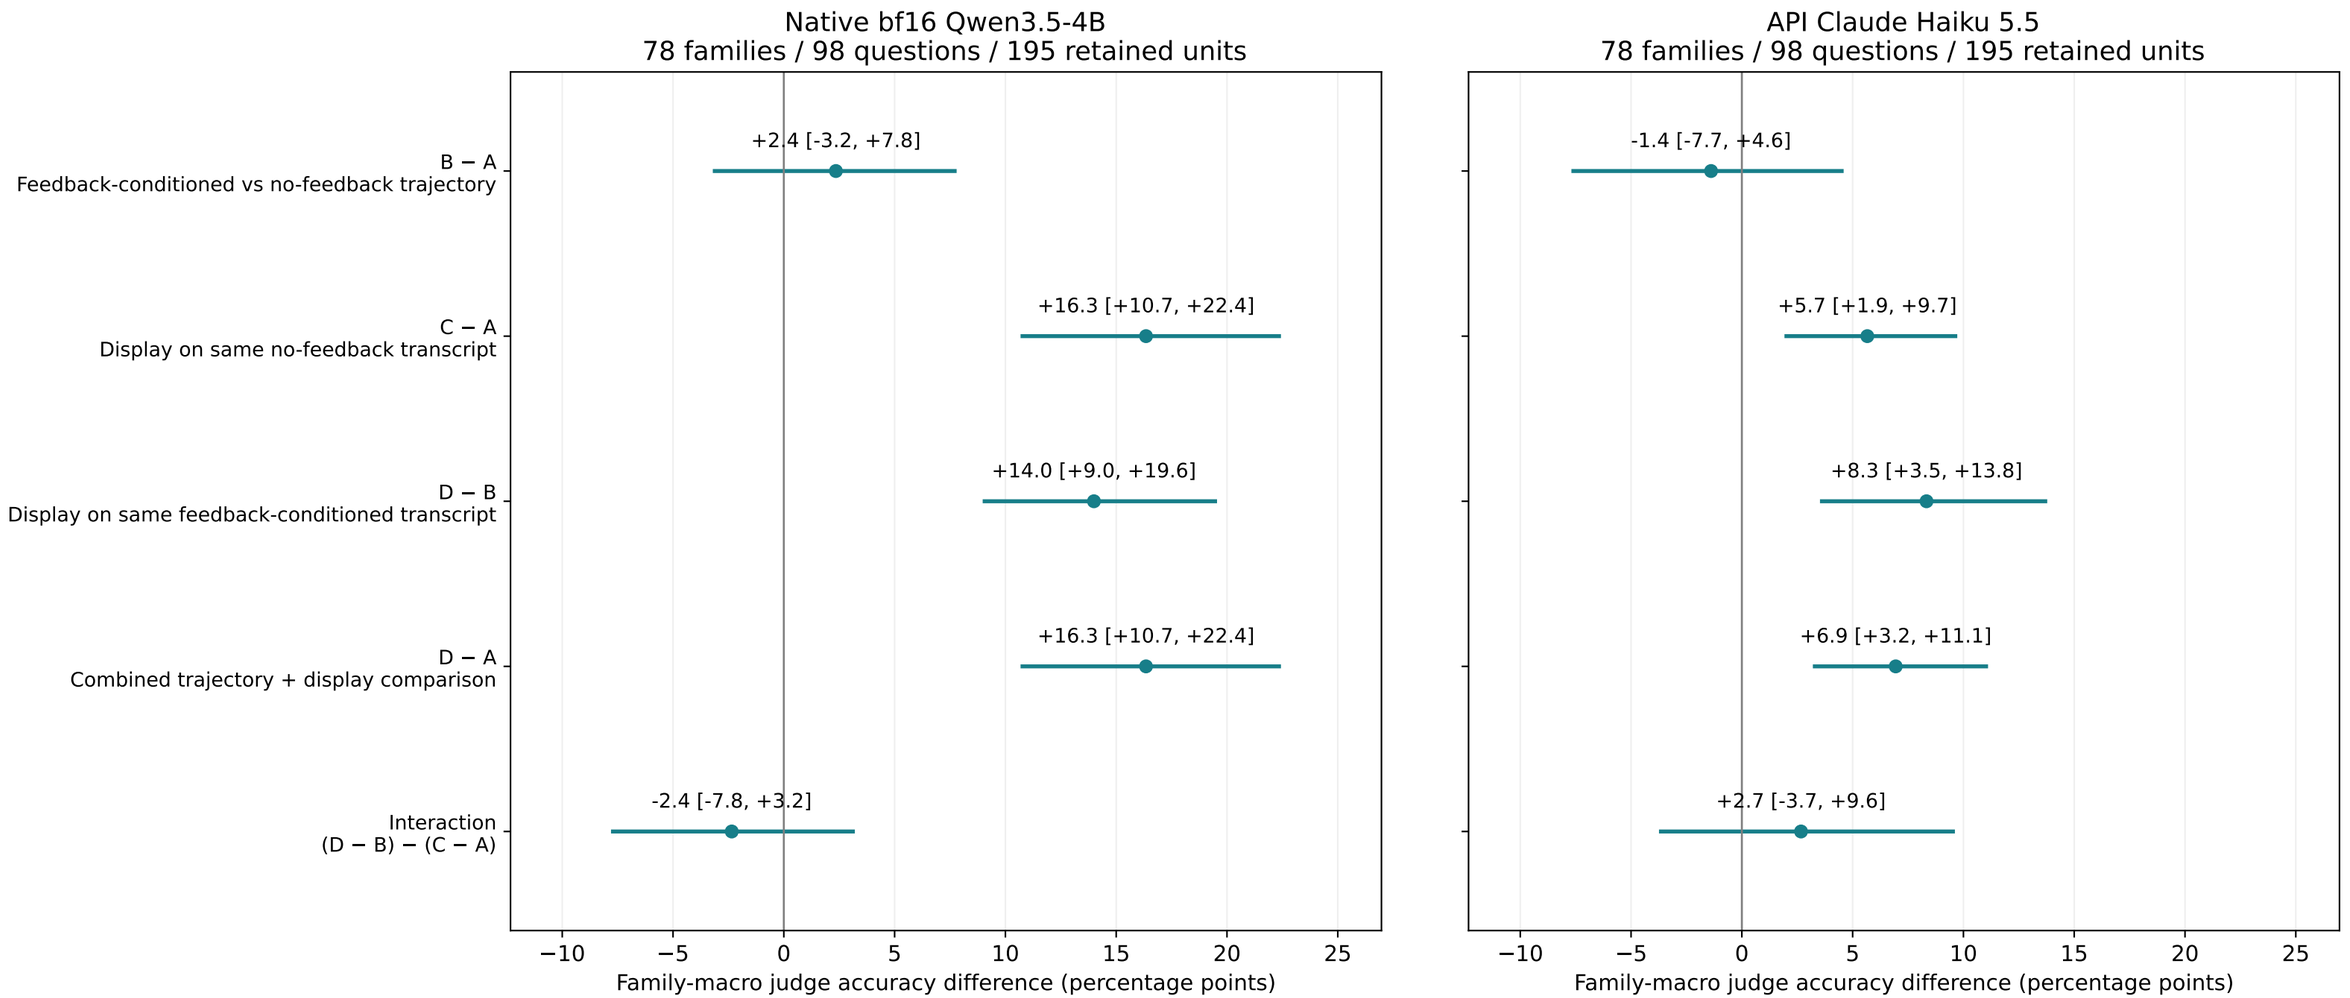

A: no feedback / no packet · B: feedback-conditioned / no packet · C: exact A transcript / packet · D: exact B transcript / packet.
Packet: frozen adapted response-mean logistic head, whole structured-response mean pooling; Qwen3.5 layer 43 / Qwen3.8 layer 35; prior-turn percentiles, speaker means, optional recommendation.
Equal-weight source-family means. Intervals: unadjusted 95% percentile family bootstrap, 10,000 resamples.
They describe family variation; one debate seed gives no repeated-generation uncertainty. B −A/D−A change trajectories; C−A/D−B hold public text fixed.
Selection note: fallible source review plus post-generation, model-specific eligibility; 195 units / 98 questions because one question retains only one model.